KNN for Loan Default Predictions

Install xlrd for Excel support and import the core analysis, plotting, and k-NN libraries.

```
In [35]: %pip install xlrd
         import pandas as pd
         import numpy as np
         import matplotlib.pyplot as plt

         from sklearn.neighbors import KNeighborsClassifier
         from sklearn.cluster import KMeans
```

Requirement already satisfied: xlrd in c:\users\betoq\onedrive\documentos\github\exc
hange-data-concepts\.venv\lib\site-packages (2.0.2)
Note: you may need to restart the kernel to use updated packages.

Load the credit default dataset from Excel, rename the target column to TARGET, and
preview the data.

```
In [36]: DATA_FILE = "default of credit card clients.xls"

         raw_df = (
             pd.read_excel(DATA_FILE, header=1)
             .rename(columns={"default payment next month": "TARGET"})
         )

         raw_df.head()
```

Out[36]:

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4	...	
0	1	20000	2	2	1	24	2	2	-1	-1	...	
1	2	120000	2	2	2	26	-1	2	0	0	...	
2	3	90000	2	2	2	34	0	0	0	0	...	
3	4	50000	2	2	1	37	0	0	0	0	...	
4	5	50000	1	2	1	57	-1	0	-1	0	...	

5 rows × 25 columns

1.1 Prepare Data

Split the dataframe into feature columns (X) and the TARGET labels (y), then display both.
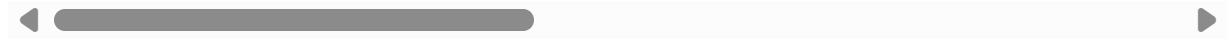
```
In [37]: X = raw_df.drop(columns = 'TARGET')

         y = raw_df['TARGET']
```

```
display(X,y)
```

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4
0	1	20000	2	2	1	24	2	2	-1	-1
1	2	120000	2	2	2	26	-1	2	0	0
2	3	90000	2	2	2	34	0	0	0	0
3	4	50000	2	2	1	37	0	0	0	0
4	5	50000	1	2	1	57	-1	0	-1	0
...	...	...	...	...	...	...	...	...	...	...
29995	29996	220000	1	3	1	39	0	0	0	0
29996	29997	150000	1	3	2	43	-1	-1	-1	-1
29997	29998	30000	1	2	2	37	4	3	2	-1
29998	29999	80000	1	3	1	41	1	-1	0	0
29999	30000	50000	1	2	1	46	0	0	0	0

30000 rows × 24 columns

```
0        1
1        1
2        0
3        0
4        0
        ..
29995    0
29996    0
29997    1
29998    1
29999    1
Name: TARGET, Length: 30000, dtype: int64
```

1.2(a) Split data

Create an 80/20 train-test split with a fixed seed and then carve out a validation fold from the training data for threshold tuning.

```python
In [38]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.2,
    random_state=2025,
    stratify=y,
)  # 20% of data is test here
```

```
X_train_sub, X_val, y_train_sub, y_val = train_test_split(
    X_train,
    y_train,
    test_size=0.25,
    random_state=0,
    stratify=y_train,
)  # carve out a validation fold (20% of original data) for threshold tuning

display(X_train_sub.shape, X_val.shape, X_test.shape, y_train_sub.shape, y_val.shap
```

```
(18000, 24)
(6000, 24)
(6000, 24)
(18000,)
(6000,)
(6000,)
```

1.2(b) Scale features

Standardize features using training-fold statistics and confirm the means and standard deviations of the scaled sets.

In [39]:
```python
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()                            # define the scaler

scaler.fit(X_train_sub)                              # train the scaler on training
X_train_scaled = scaler.transform(X_train_sub)       # apply the scaler to transform
X_val_scaled = scaler.transform(X_val)               # transform validation data
X_test_scaled = scaler.transform(X_test)             # apply the scaler to transform

dataframe_scaled = pd.DataFrame(np.round(X_train_scaled, 2), columns=X.columns)
dataframe_scaled
dataframe_scaled.describe()
```

Out[39]:

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE
count	18000.000000	18000.000000	18000.000000	18000.000000	18000.000000	18000.000000
mean	0.000010	0.000136	-0.000319	-0.001118	0.000288	-0.000084
std	1.000009	1.000311	1.002294	0.999471	0.998925	0.999967
min	-1.740000	-1.220000	-1.240000	-2.340000	-2.960000	-1.570000
25%	-0.870000	-0.910000	-1.240000	-1.080000	-1.050000	-0.820000
50%	0.000000	-0.210000	0.810000	0.180000	0.860000	-0.170000
75%	0.860000	0.570000	0.810000	0.180000	0.860000	0.690000
max	1.730000	4.910000	0.810000	5.230000	2.770000	4.680000

8 rows × 24 columns

◀ ▬▬▬▬▬▬▬▬▬▬ ▶

1.3 Modeling

1.3(a) find best K

Set up a k-NN classifier and run a 5-fold grid search over odd k values to find the best parameter.

In [40]:
```python
from sklearn.model_selection import GridSearchCV

knn = KNeighborsClassifier()

k_range = {'n_neighbors': np.arange(1,80,2)}                          #  k = 1,3,

grid = GridSearchCV(estimator = knn, param_grid = k_range, cv = 5)      # 5-CV on t

grid.fit(X_train_scaled, y_train_sub)        # search over the values on Train data

best_param = grid.best_params_              # k value that returns highest mean cv scor
best_cv_score = grid.best_score_            # mean cv score of the best k

print("Best Params: {}".format(best_param))
print("Mean cv score of the best k: {:.2%}".format(best_cv_score))
```
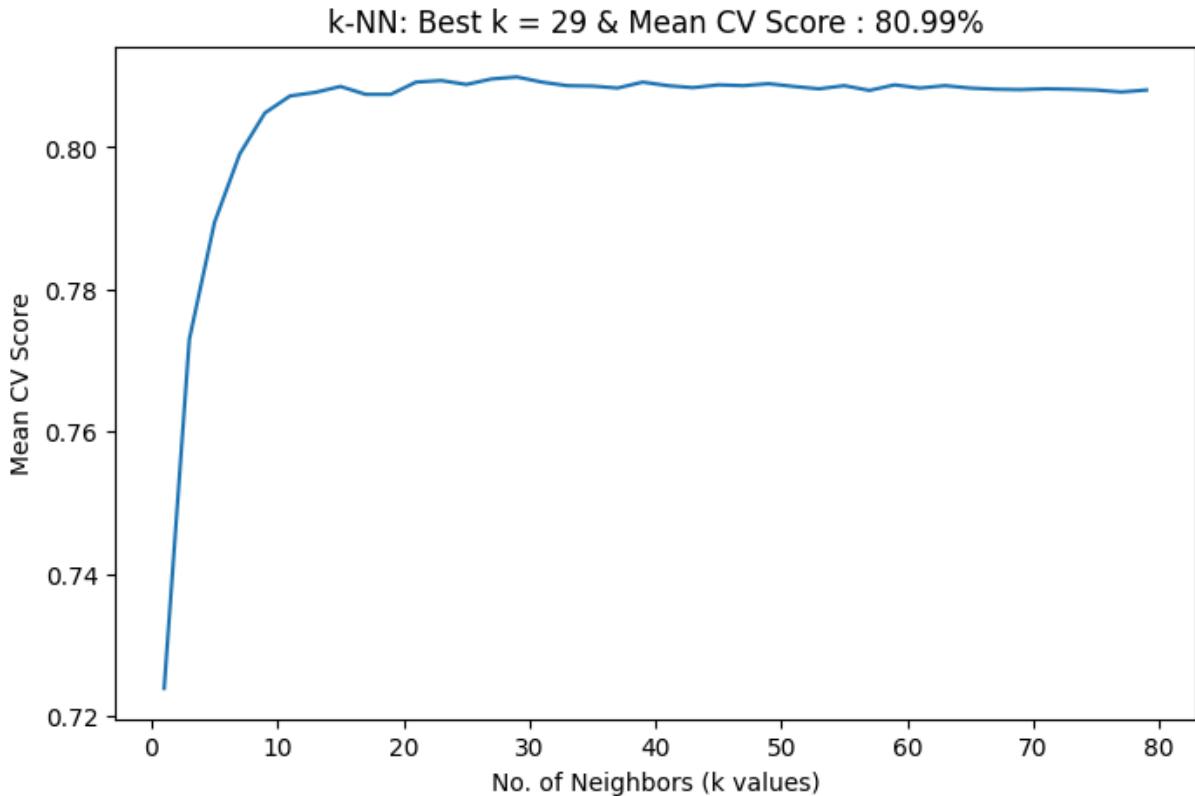
```
Best Params: {'n_neighbors': np.int64(29)}
Mean cv score of the best k: 80.99%
```

1.3(b) Visualize mean cv scores for each k value

Plot mean cross-validation accuracy versus k to visualize how performance changes with neighborhood size.

```
In [41]: k_values = grid.cv_results_['param_n_neighbors']          # k values, same as k_range
         cv_scores = grid.cv_results_['mean_test_score']           # mean cv scores for each k

         plt.figure(figsize = (8, 5))
         plt.plot(k_values,  cv_scores)         # a line plot (default blue solid line)
         plt.xlabel('No. of Neighbors (k values)')
         plt.ylabel('Mean CV Score')
         plt.title('k-NN: Best k = {} & Mean CV Score : {:.2%}'.format(best_param['n_neighbo
         plt.show()
```


k-NN: Best k = 29 & Mean CV Score : 80.99%

Compute average credit limits for defaulters and non-defaulters to ground the cost assumptions.

```
In [42]: # 1. Separate defaulters vs non-defaulters
         df_default    = raw_df[raw_df["TARGET"] == 1]  # will default
         df_no_default = raw_df[raw_df["TARGET"] == 0]  # will NOT default

         # 2. Typical credit limit (principal proxy)
         mean_limit_default    = df_default["LIMIT_BAL"].mean()
         mean_limit_no_default = df_no_default["LIMIT_BAL"].mean()

         print("Avg credit limit (default):    ", mean_limit_default)
         print("Avg credit limit (no default):", mean_limit_no_default)
```

```
Avg credit limit (default):     130109.65641952984
Avg credit limit (no default): 178099.72607430234
```

Apply the PDF's six-month observation window and APR assumptions to derive per-period profit/loss inputs for the cost matrix.

```python
In [43]:   # Business assumptions grounded in the PDF's six-month observation window
           # April–September 2005 appears in the data dictionary PDF, so model six monthly bil
           assumption_config = {
               "annual_apr": 0.18,              # 18% annual percentage rate from the reference
               "periods_per_year": 12,          # monthly compounding
               "observation_months": 6,         # six billing cycles in the PDF window (Apr–Sep
               "loss_given_default": 0.5,       # lose 50% of the limit if they default
           }

           annual_apr = assumption_config["annual_apr"]
           periods_per_year = assumption_config["periods_per_year"]
           observation_months = assumption_config["observation_months"]
           loss_given_default = assumption_config["loss_given_default"]

           periodic_rate = annual_apr / periods_per_year
           period_length_months = 12 / periods_per_year
           periods_in_window = observation_months / period_length_months

           # Approx profit from approving a good customer (non-defaulter) over the observation
           profit_good = mean_limit_no_default * ((1 + periodic_rate) ** periods_in_window - 1

           # Approx loss from approving a bad customer (defaulter)
           loss_bad = mean_limit_default * loss_given_default

           print("Periodic rate: {:.2%}".format(periodic_rate))
           print("Approx profit per good customer: {:.2f}".format(profit_good))
           print("Approx loss per bad customer:    {:.2f}".format(loss_bad))
```

```
Periodic rate: 1.50%
Approx profit per good customer: 16642.22
Approx loss per bad customer:     65054.83
```

Construct a cost/benefit matrix that maps actual/predicted outcomes to monetary values using the assumptions.

```python
In [44]:   # 3. Cost/benefit matrix driven by the data

           # Story:
           # TARGET=1 = will default
           # Prediction=1 = treat as risky (reject / restrict credit)
           # Prediction=0 = treat as safe (approve / keep credit)

           value_TN =  profit_good     # Actual 0, Pred 0: good customer, approved → earn intere
           value_FP =  0.0             # Actual 0, Pred 1: good customer, rejected → lose that p
           value_FN = -loss_bad        # Actual 1, Pred 0: bad customer, approved → lose money
           value_TP =  0.0             # Actual 1, Pred 1: bad customer, rejected → avoided loss

           value_matrix = pd.DataFrame(
               {
                   0: {0: value_TN, 1: value_FN},   # Predicted 0
                   1: {0: value_FP, 1: value_TP},   # Predicted 1
               }
           )
           value_matrix.index.name = "Actual"
           value_matrix.columns.name = "Predicted"
```

```
value_matrix
```

Out[44]:

Predicted	0	1
Actual		
0	16642.219712	0.0
1	-65054.828210	0.0

Define a helper that summarizes classification metrics, rates, AUCs, and optional utility for a model.

In [45]:

```python
from sklearn.metrics import (
    average_precision_score,
    confusion_matrix,
    f1_score,
    precision_score,
    recall_score,
    roc_auc_score,
)

import numpy as np
import pandas as pd


def summarize_model(
    model_name,
    dataset_name,
    y_true,
    y_pred,
    y_score,
    threshold,
    utility_fn=None,
):
    '''Return a one-row DataFrame with key classification metrics.'''
    y_true_arr = np.asarray(y_true).ravel()
    y_pred_arr = np.asarray(y_pred).ravel()

    score_arr = None
    if y_score is not None:
        score_arr = np.asarray(y_score)
        if score_arr.ndim > 1:
            score_arr = score_arr[:, -1]
        score_arr = score_arr.ravel()

    tn, fp, fn, tp = confusion_matrix(
        y_true_arr, y_pred_arr, labels=[0, 1]
    ).ravel()
    total = tn + fp + fn + tp

    def _safe_div(num, denom):
        return num / denom if denom else np.nan
```

```python
        fprate = _safe_div(fp, fp + tn)
        fnrate = _safe_div(fn, fn + tp)
        approval_rate = _safe_div(tn + fn, total)
        rejection_rate = _safe_div(fp + tp, total)

        roc_auc = np.nan
        pr_auc = np.nan
        if score_arr is not None and len(np.unique(y_true_arr)) > 1:
            try:
                roc_auc = roc_auc_score(y_true_arr, score_arr)
            except ValueError:
                pass
            try:
                pr_auc = average_precision_score(y_true_arr, score_arr)
            except ValueError:
                pass

        precision = precision_score(y_true_arr, y_pred_arr, zero_division=0)
        recall = recall_score(y_true_arr, y_pred_arr, zero_division=0)
        f1 = f1_score(y_true_arr, y_pred_arr, zero_division=0)

        utility_total = np.nan
        utility_per_app = np.nan
        if utility_fn is not None:
            utility_result = utility_fn(tn, fp, fn, tp)
            if isinstance(utility_result, dict):
                utility_total = utility_result.get('utility_total', np.nan)
                utility_per_app = utility_result.get('utility_per_app', np.nan)
            elif isinstance(utility_result, (tuple, list)):
                if len(utility_result) >= 2:
                    utility_total, utility_per_app = utility_result[:2]
                elif len(utility_result) == 1:
                    utility_total = utility_result[0]
            else:
                utility_total = utility_result

            try:
                needs_per_app = np.isnan(utility_per_app)
            except TypeError:
                needs_per_app = False

            if needs_per_app:
                utility_per_app = _safe_div(utility_total, total)

        summary = {
            'model_name': model_name,
            'dataset_name': dataset_name,
            'threshold': threshold,
            'TN': tn,
            'FP': fp,
            'FN': fn,
            'TP': tp,
            'fprate': fprate,
            'fnrate': fnrate,
            'approval_rate': approval_rate,
            'rejection_rate': rejection_rate,
```

```
            'roc_auc': roc_auc,
            'pr_auc': pr_auc,
            'precision': precision,
            'recall': recall,
            'f1': f1,
            'utility_total': utility_total,
            'utility_per_app': utility_per_app,
        }

        return pd.DataFrame([summary])
```

Find the probability threshold that maximizes expected utility using validation data and summarize the tuned model on the test data.

In [46]:
```python
# Summarize the tuned k-NN model using a validation-tuned threshold on the held-out

def _knn_utility(tn, fp, fn, tp):
    total = (tn * value_TN) + (fp * value_FP) + (fn * value_FN) + (tp * value_TP)
    total_apps = tn + fp + fn + tp
    per_app = total / total_apps if total_apps else np.nan
    return {
        'utility_total': total,
        'utility_per_app': per_app,
    }


y_score_val = grid.predict_proba(X_val_scaled)[:, 1]
_knn_model_name = "k-NN (k={})".format(grid.best_params_['n_neighbors'])

threshold_grid = np.linspace(0.0, 1.0, 101)
threshold_summaries = []
for _thr in threshold_grid:
    y_pred_thr = (y_score_val >= _thr).astype(int)
    threshold_summaries.append(
        summarize_model(
            model_name=_knn_model_name,
            dataset_name="Validation",
            y_true=y_val,
            y_pred=y_pred_thr,
            y_score=y_score_val,
            threshold=_thr,
            utility_fn=_knn_utility,
        )
    )

threshold_results = pd.concat(threshold_summaries, ignore_index=True)

best_idx = threshold_results['utility_per_app'].fillna(-np.inf).idxmax()
best_threshold = threshold_results.loc[best_idx, 'threshold']

y_score_test = grid.predict_proba(X_test_scaled)[:, 1]
y_pred_best = (y_score_test >= best_threshold).astype(int)
knn_summary_df = summarize_model(
    model_name=_knn_model_name,
    dataset_name="Test",
```

```
        y_true=y_test,
        y_pred=y_pred_best,
        y_score=y_score_test,
        threshold=best_threshold,
        utility_fn=_knn_utility,
    )

    print(best_threshold)

    knn_summary_df
```

0.21

Out[46]:

	model_name	dataset_name	threshold	TN	FP	FN	TP	fprate	fnrate	appro
0	k-NN (k=29)	Test	0.21	3550	1123	532	795	0.240317	0.400904	

◄ ▬▬▬▬▬▬▬▬▬▬▬▬▬▬▬▬▬ ►

1.3(c) Evaluate and apply the best model (validation-tuned)

Use the validation-tuned threshold to generate test-set predictions and view the label counts.

In [47]:
```
y_pred_test = y_pred_best

display(y_test, y_pred_test)
```

```
17479    0
11591    0
23496    0
6564     0
15547    0
        ..
20838    1
22958    0
17594    0
24251    0
1942     0
Name: TARGET, Length: 6000, dtype: int64
array([0, 1, 0, ..., 1, 0, 0], shape=(6000,))
```

2. Model Evaluation

Confusion matrix based on the tuned threshold.
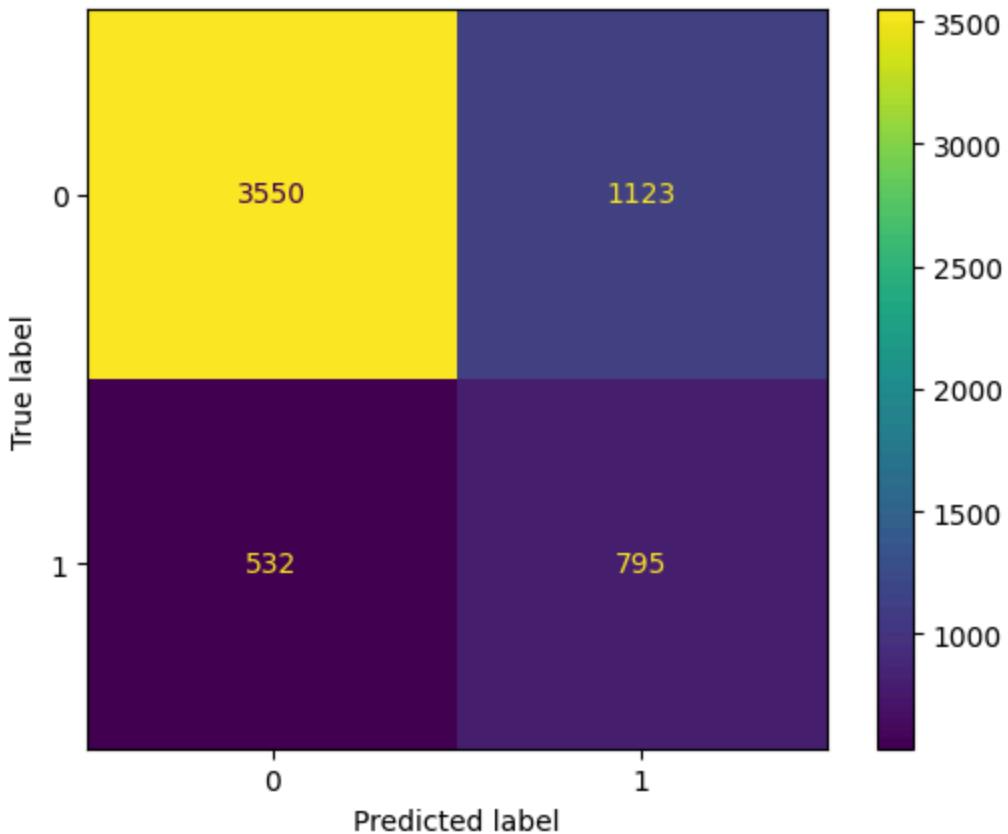
In [48]:
```
from sklearn.metrics import ConfusionMatrixDisplay

disp = ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test, display_labels=
```

Tabulate the raw confusion matrix counts for the test set using the tuned cutoff.

```
In [49]:  cm_counts = pd.crosstab(
              y_test,
              y_pred_test,
              rownames=['Actual'],
              colnames=['Predicted'],
              dropna=False
          )

          cm_counts
```

Out[49]:

Predicted	0	1
Actual		
0	3550	1123
1	532	795

Convert confusion matrix counts into probabilities for expected-value calculations driven by the tuned model.

```
In [50]:  # 2.2 Convert confusion matrix counts to probabilities (expected rates)

          N = cm_counts.to_numpy().sum()        # total number of test samples
          cm_prob = cm_counts / N               # element-wise division
```

cm_prob

Out[50]:

Predicted	0	1
Actual		
0	0.591667	0.187167
1	0.088667	0.132500

Calculate expected value per applicant and for the full test set using the tuned confusion matrix.

In [51]:
```python
# 2.4 Expected value per person and in total

# EV per person = sum over all cells of [probability * value]
ev_per_person = (cm_prob * value_matrix).to_numpy().sum()

# EV over all people in the test set (just multiply by N)
ev_total = ev_per_person * N

print("Expected value per person: {:.2f}".format(ev_per_person))
print("Expected value over all {} test customers: {:.2f}".format(N, ev_total))
```

```
Expected value per person: 4078.45
Expected value over all 6000 test customers: 24470711.37
```