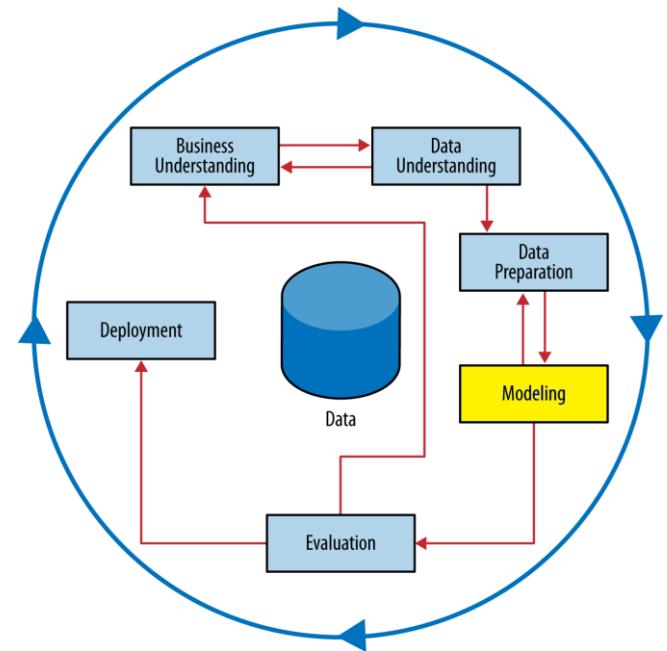


Predictive Modeling: Decision Tree

PF3

Decision Tree and Predictive Modeling

- ▶ Predictive Modeling: a Classification Example
 - ▶ Model training: tree induction
 - ▶ Supervised learning
 - ▶ Predict with a classification tree
 - ▶ Class prediction
 - ▶ Class probability estimation
- ▶ Decision Tree
 - ▶ Identify **informative attributes** to predict the **target** value
 - ▶ Segment data by **recursive attribute selection**



Supervised Segmentation with Tree

- ▶ The **target** attribute supervises the model training process.
 - ▶ Instances (e.g., customers) are segmented into subgroups according to their **feature** values.
 - ▶ The goal is to have members in each subgroup with similar **target** values.
- ▶ **Informative attributes**
 - ▶ A **feature** is **informative** if it can help us learn more about the **target** attribute.

This is one row (example).
Feature vector is: <Claudio,115000,40,no>
Class label (value of Target attribute) is **no**

Attributes

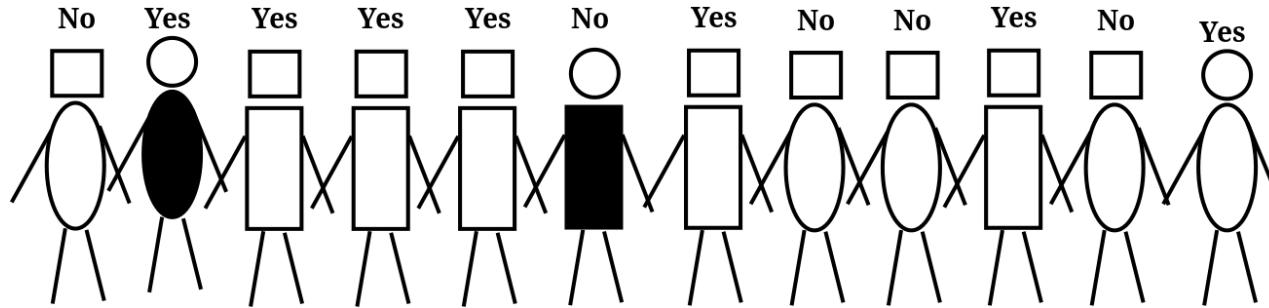
Name	Balance	Age	Employed	Write-off
Mike	\$200,000	42	no	yes
Mary	\$35,000	33	yes	no
Claudio	\$115,000	40	no	no
Robert	\$29,000	23	yes	yes
Dora	\$72,000	31	no	no

Features

Target Attribute

Select Informative Attributes

- ▶ **Objective:** based on customer **features**, segment them into subgroups so that we can best distinguish **write-offs** (reject) from **no write-offs** (approve).
- ▶ In each subgroup, have as many instances with the **same target value** (i.e., class label) as possible.



Features

Head-shape: Square, Circle

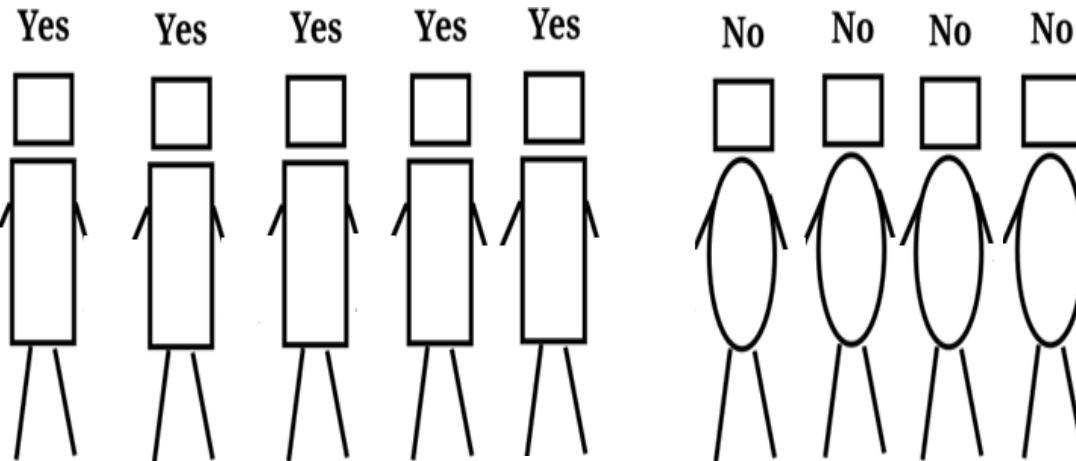
Body-shape: Rectangular, Oval

Body-Color: Black, White

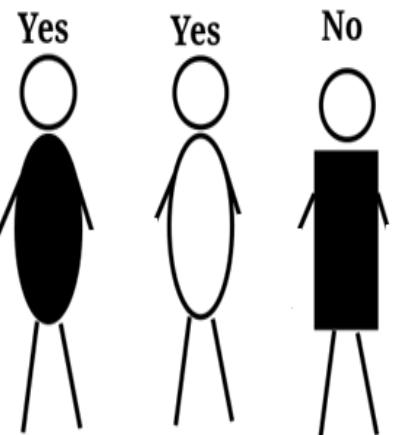
Target

Loan Write-off: Yes, No

Segment by Head-Shape?



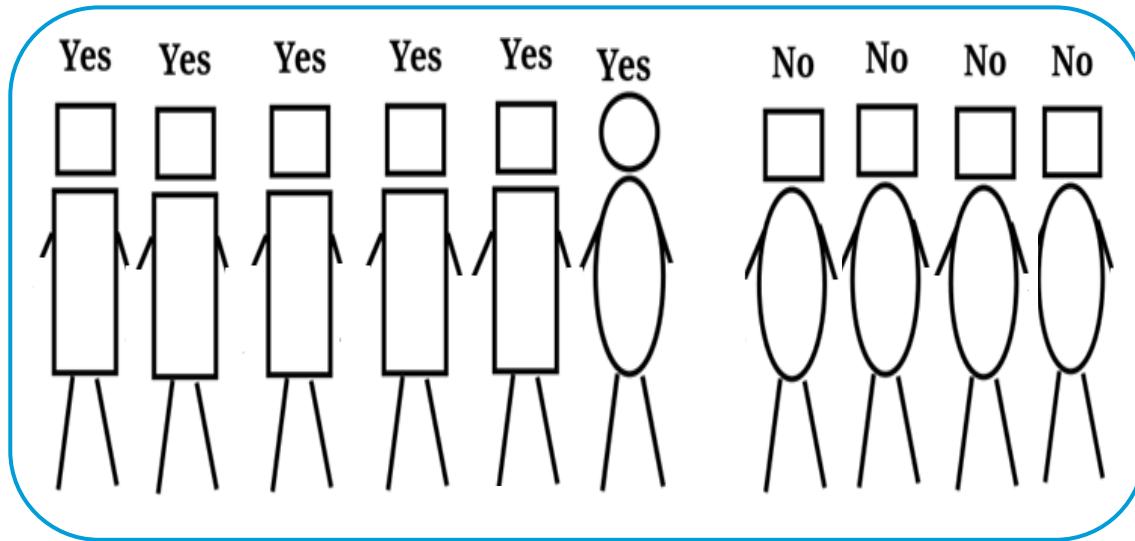
Group 1 (**Square Head**)
Result: 5 Yes 4 No



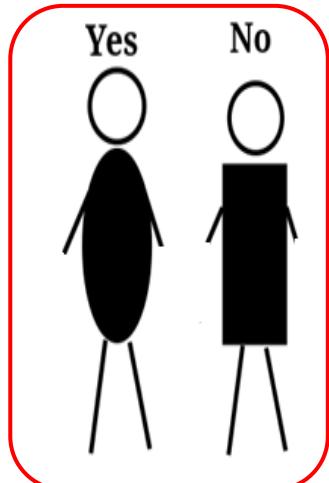
Group 2 (**Circle Head**)
Result: 2 Yes 1 No

- Very impure in both groups
- End the segmentation? Still room for improvement as the feature values are not homogenous.
 - Group 1 (Rectangular vs Oval body)
 - Group 2 (Oval vs Rectangular body)

Segment by Body-Color?



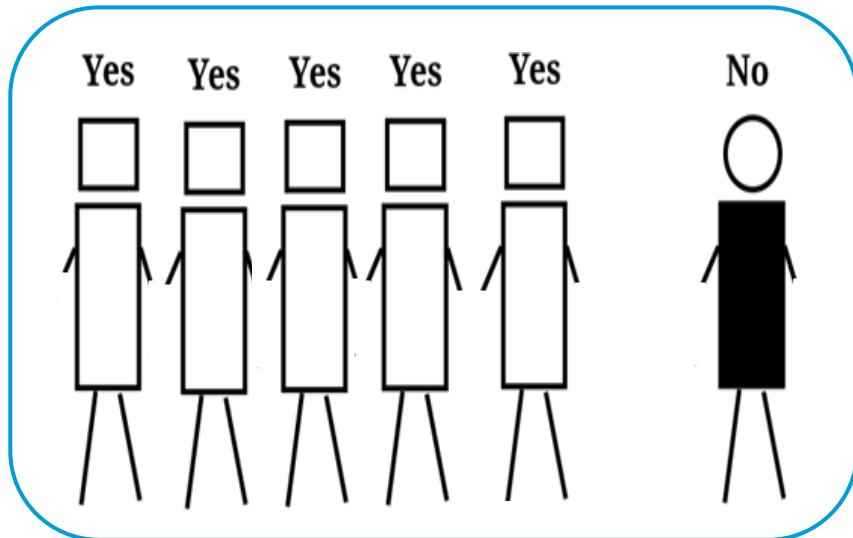
Group 1 (White Body)
Result: 6 Yes 4 No



Group 2 (Black Body)
Result: 1 Yes 1 No

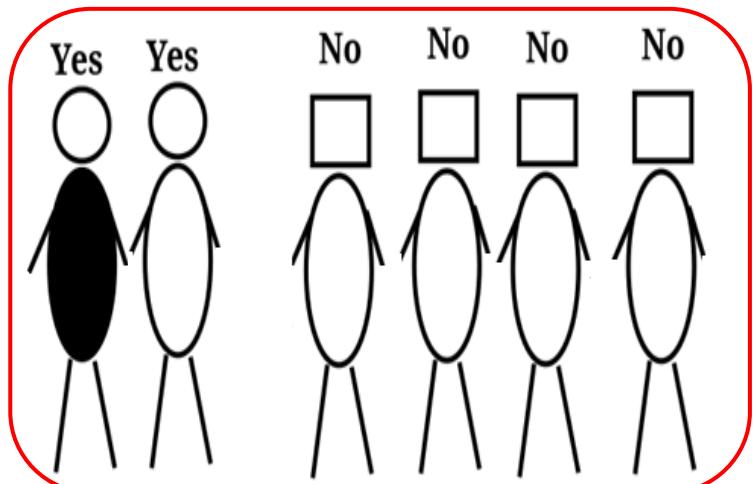
- Very impure in both groups.
- End segmentation? Still room for improvement.
 - Group 1 (Rectangular vs Oval body)
 - Group 2 (Rectangular vs Oval body)

Segment by Body-Shape?



Group 1 (**Rectangular Body**)
Result: 5 Yes 1 No

- Purer than the earlier two.
- End segmentation? Still room for improvement.
 - Group 1 (White vs Black body color)
 - Group 2 (Circle vs Square head)

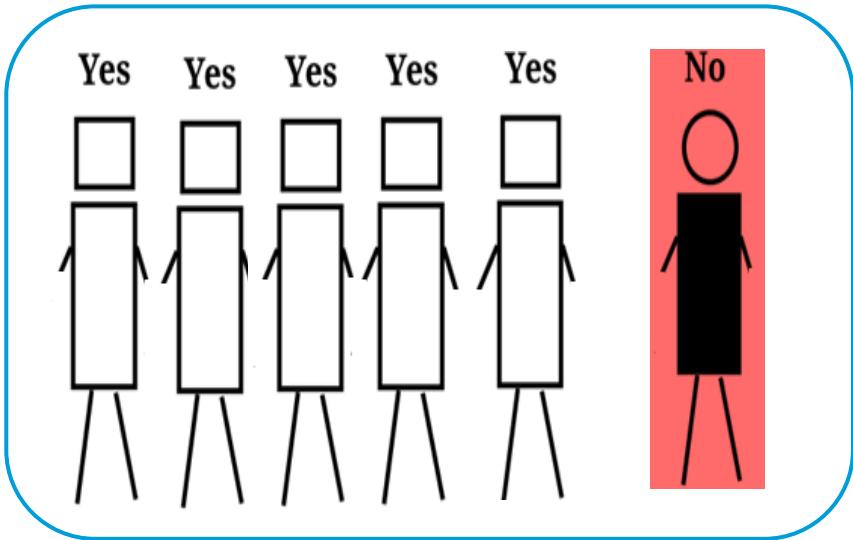


Group 2 (**Oval Body**)
Result: 2 Yes 4 No

What's Next?

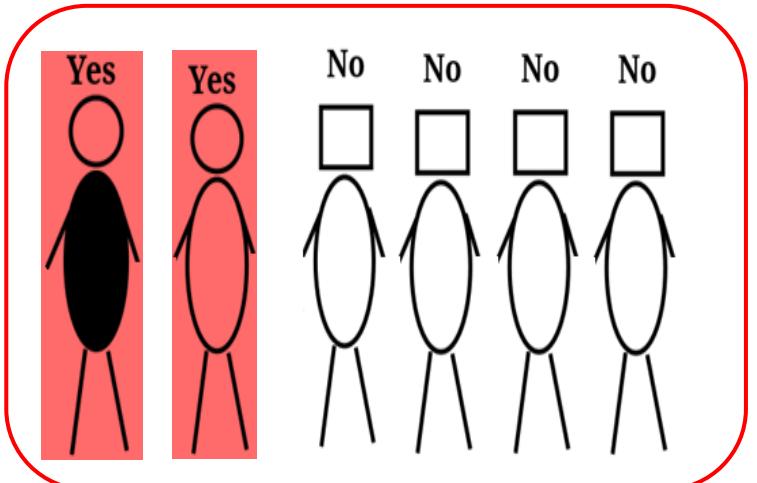
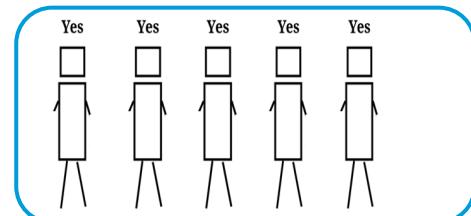
- ▶ **Body-shape** seems to be the best feature in segmenting customers in **the first split**.
 - ▶ How to quantify “the usefulness” of each feature?
- ▶ After the first split with **body-shape**, each subgroup is still not pure in terms of the target. So further segmentation for each subgroup?
 - ▶ Which feature should be used in the second split?

Further Segmentation (Second Split)



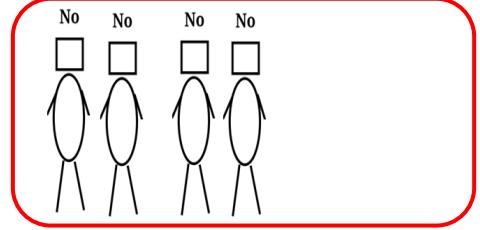
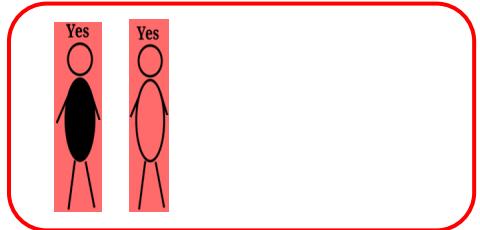
Group I
(Rectangular Body)
Result: 5 Yes 1 No

Further segment
Group I by
body color

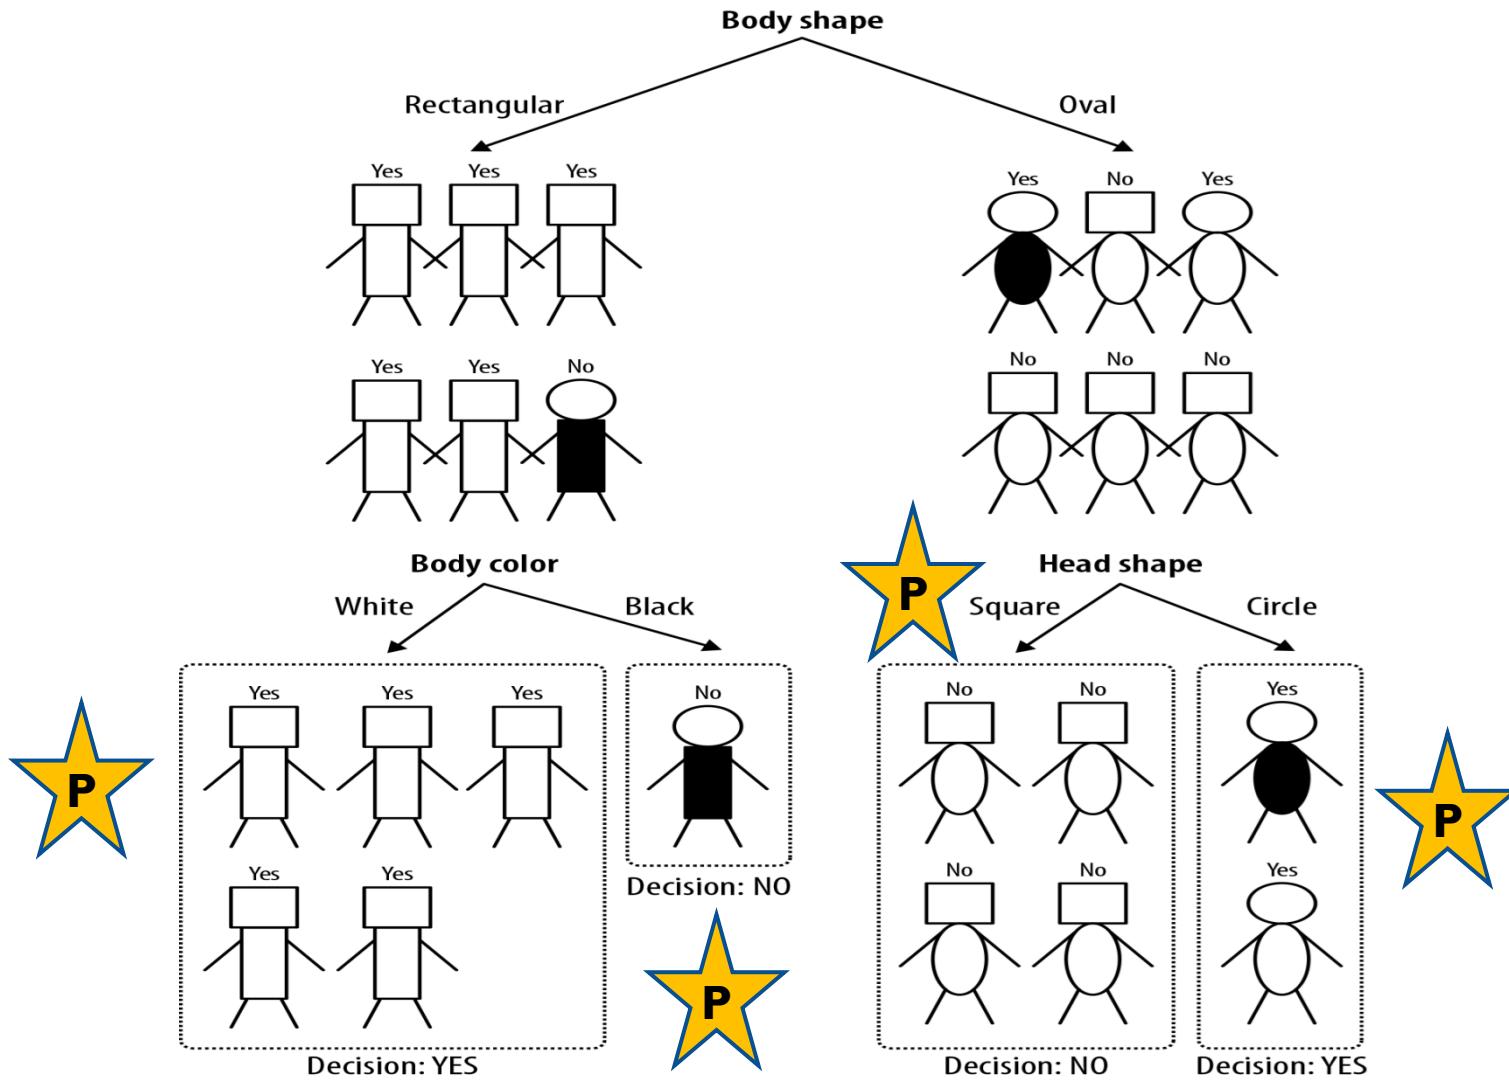


Group 2 (Oval Body)
Result: 2 Yes 4 No

Further segment
Group 2 by
head shape



Final Model



Informative Attributes

- ▶ Subgroups are expected to be as **pure** as possible.
 - ▶ A pure group: all members in a group have the same target value.
 - ▶ An impure group: at least one member has a different target value.
- ▶ **Entropy** is a measure of disorder (i.e., impurity) in a group.

$$\text{Entropy} \downarrow \Rightarrow \text{Purity} \uparrow$$

- ▶ Which feature is the best in splitting a node? We compare their **information gain** in a split.
 - ▶ **Information gain (IG)** captures the **change of entropy** (impurity) in the data, due to the split with a feature.
 - ▶ If a feature can **reduce the impurity** (entropy) substantively, then it is **informative** in predicting the target.

Entropy

- ▶ Entropy is computed with:

$$\text{entropy} = -p_1 \log_2(p_1) - p_2 \log_2(p_2) - \cdots - p_n \log_2(p_n)$$

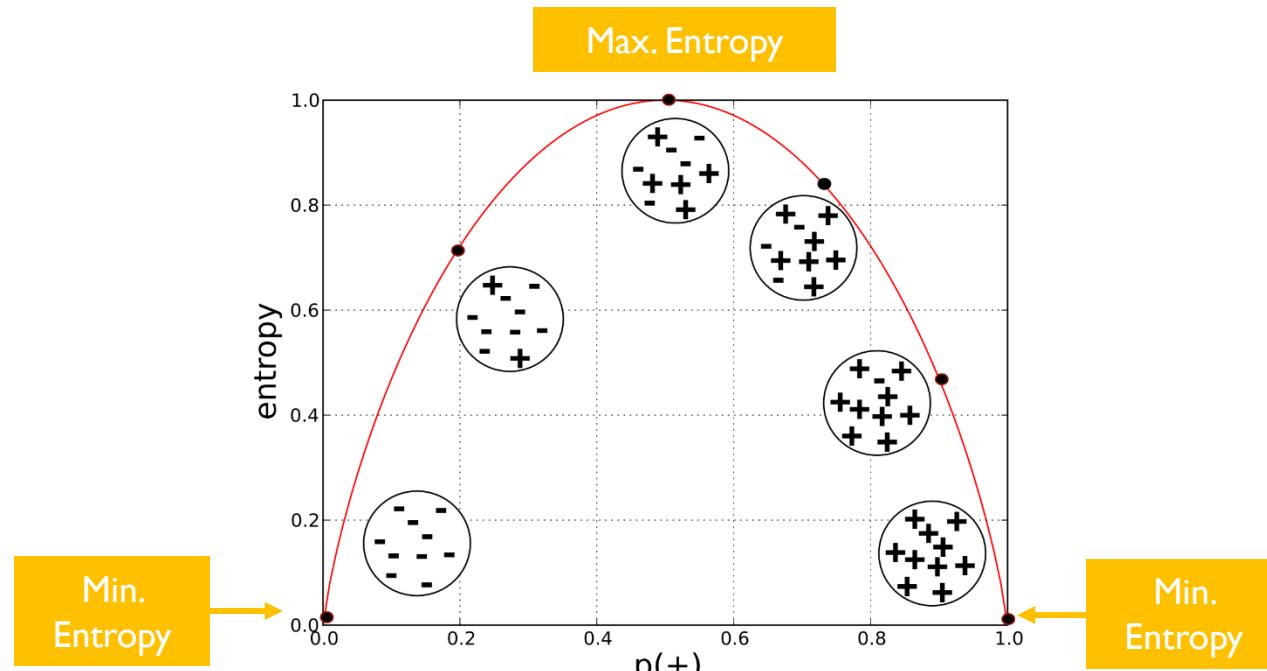
- ▶ n is the number of classes for the target
 - ▶ e.g., $n = 2$ for a binary target attribute
 - ▶ p_i is the proportion of instances of the i^{th} class in the group.
 - ▶ \log_2 is logarithm with base 2
-
- ▶ Entropy is ranging from
 - ▶ 0 at **minimum disorder** (all members the same), to
 - ▶ 1 at **maximum disorder** (classes are equally mixed)

Entropy for a Two-class Group

- In a binary classification task where the target attribute has only two classes ($n = 2$): + and -.

$$\text{entropy} = -p_1 \log_2(p_1) - p_2 \log_2(p_2)$$

- p_1 and p_2 are the proportions of + and - class in a group



Entropy for Parent and Child Nodes

Parent

No segmentation yet

Entropy(p)

$$(-7/12) \cdot \log_2(7/12) - (5/12) \cdot \log_2(5/12) = 0.98$$

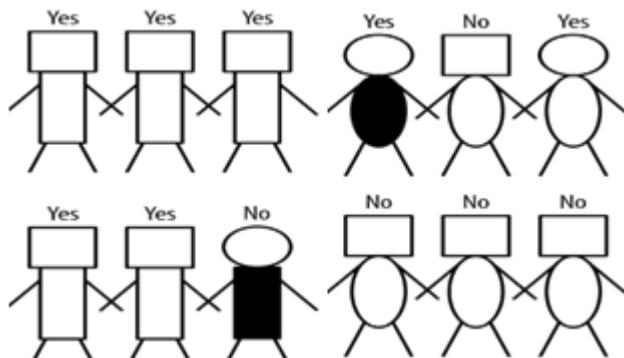
2 Children

One feature is used for splitting.

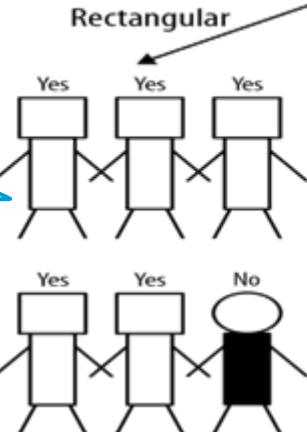
Entropy(c)

$$\text{Entropy}(c_1) = 0.65$$

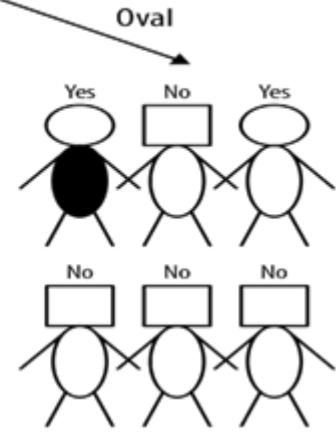
$$\text{Entropy}(c_2) = 0.92$$



Body shape



Child 1



Child 2

Information Gain

- ▶ **Information gain** measures the change of entropy due to the use of any feature in a split.

$$\text{Information Gain} = \text{entropy}(\text{parent}) - \text{entropy}(\text{children})$$

- ▶ Each child node (c_i) contains $p(c_i)$ of the total instances.
 - ▶ If there are m children after splitting
- ▶ **Entropy for all children:** sum of all child nodes' entropy weighted by the proportion of instances in each node.

$$\text{entropy}(\text{children}) = \sum_{i=1}^m p(c_i) \times \text{entropy}(c_i)$$

Information Gain of *Body Shape*

Entropy(p)

$$(-7/12)\log_2(7/12) - (5/12)\log_2(5/12) = 0.98$$

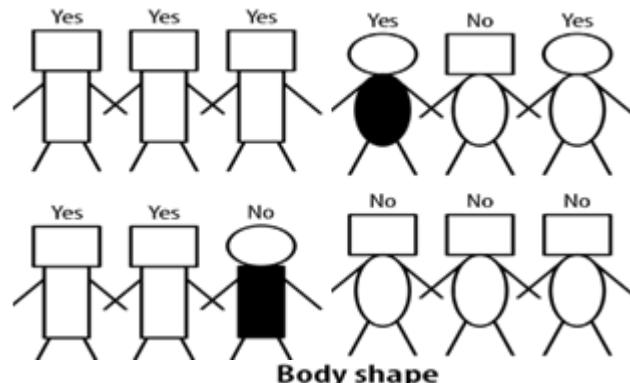
$IG(\text{body shape}) = 0.2$

Entropy(c)

$$\text{Entropy}(c_1) = 0.65$$

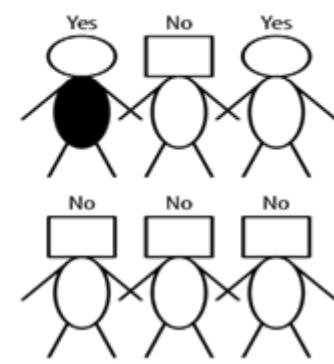
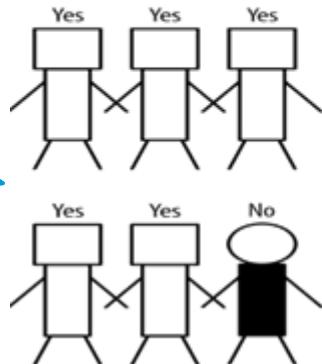
$$\text{Entropy}(c_2) = 0.92$$

$$\text{Entropy}(c) = (1/2) * 0.65 + (1/2) * 0.92 = 0.78$$



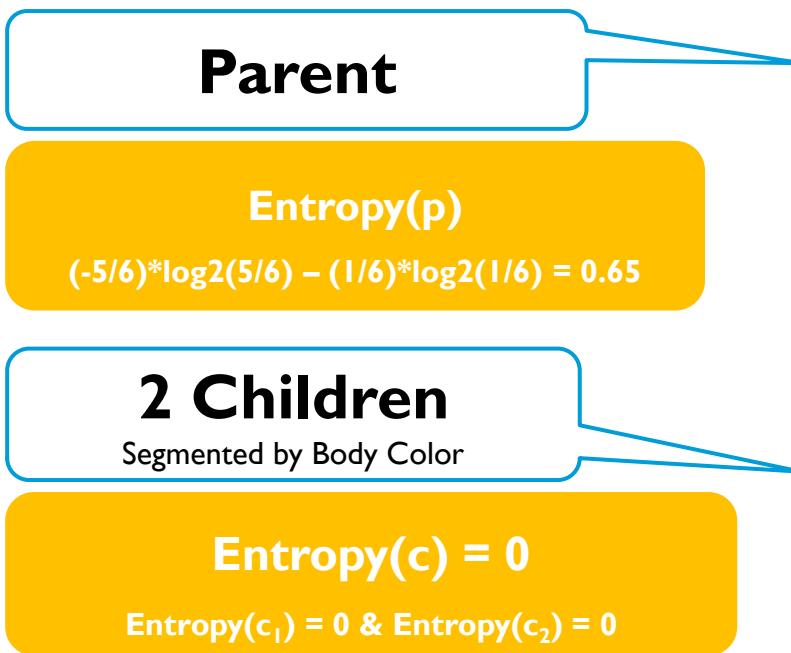
Rectangular

Oval

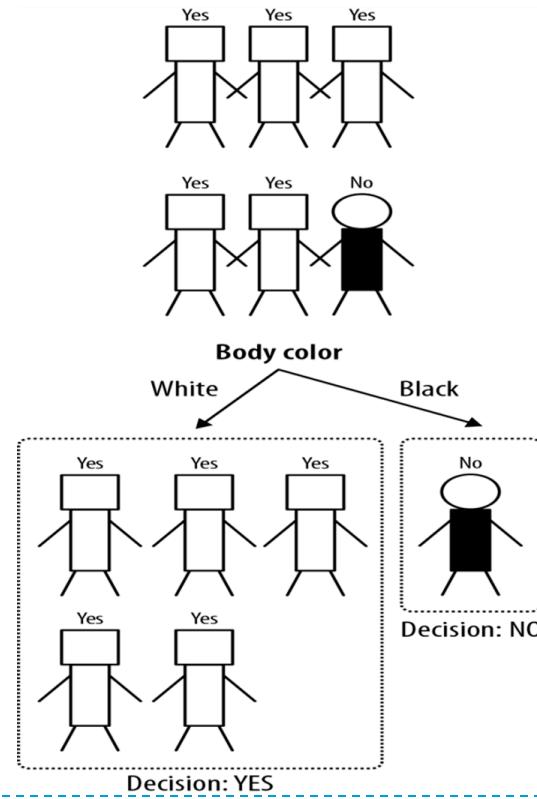


Recursive Feature Selection

- ▶ The child node in the first split (e.g., rectangular body shape) now become the parent node in second splitting.
- ▶ Need to calculate and compare the information gain for features **recursively** in splitting each subgroup.



$$IG (\text{Body Color}) = 0.65$$



Compute Information Gain for *Balance*

▶ Target:

- ▶ write-off (\star) and not write-off (\bullet)

▶ Feature:

- ▶ Outstanding balance ($< 50K$ or otherwise)
 - ▶ $c_1 = \text{Balance} < 50K$
 - ▶ $c_2 = \text{Balance} \geq 50K$

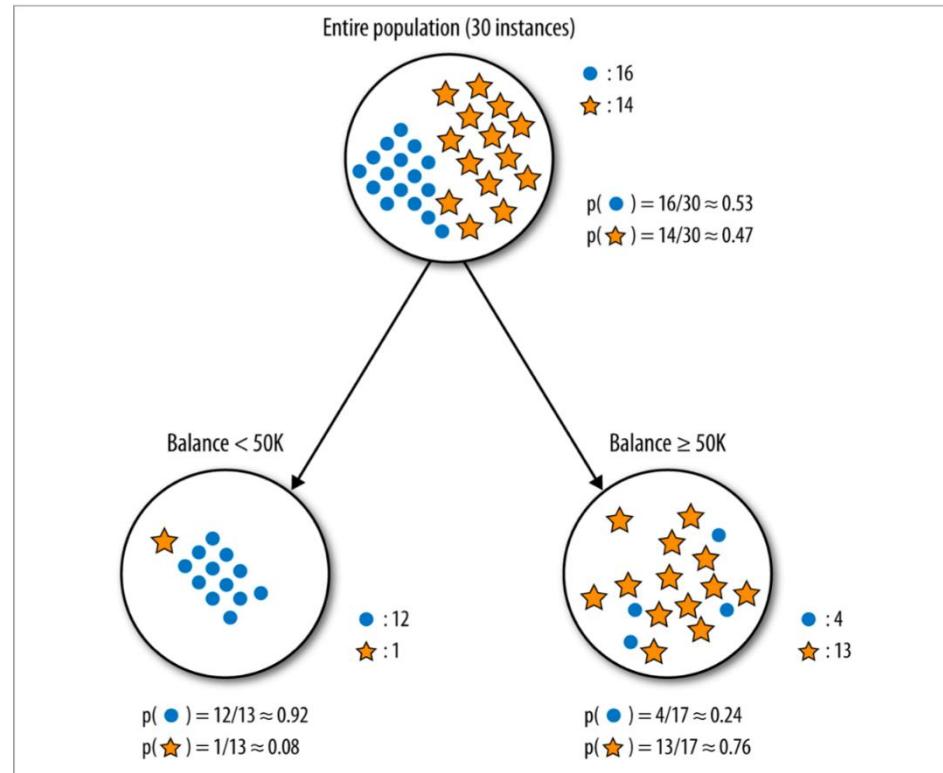
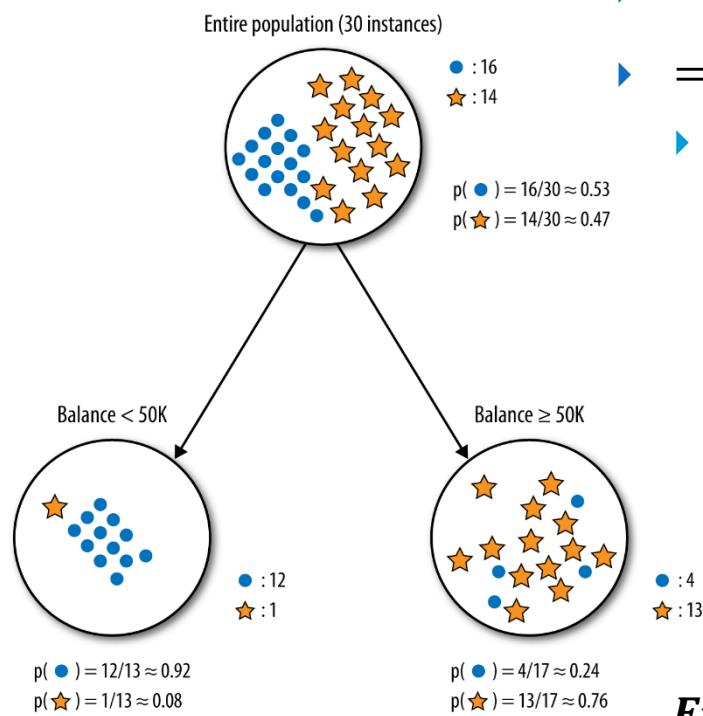


Figure 3-4. Splitting the “write-off” sample into two segments, based on splitting the Balance attribute (account balance) at 50K.

Step 1 – Entropy of Parent & Child Nodes



Entropy(c_1)

- $= -\frac{12}{13} \times \log_2 \left(\frac{12}{13} \right) - \frac{1}{13} \times \log_2 \left(\frac{1}{13} \right)$
- $= 0.39$

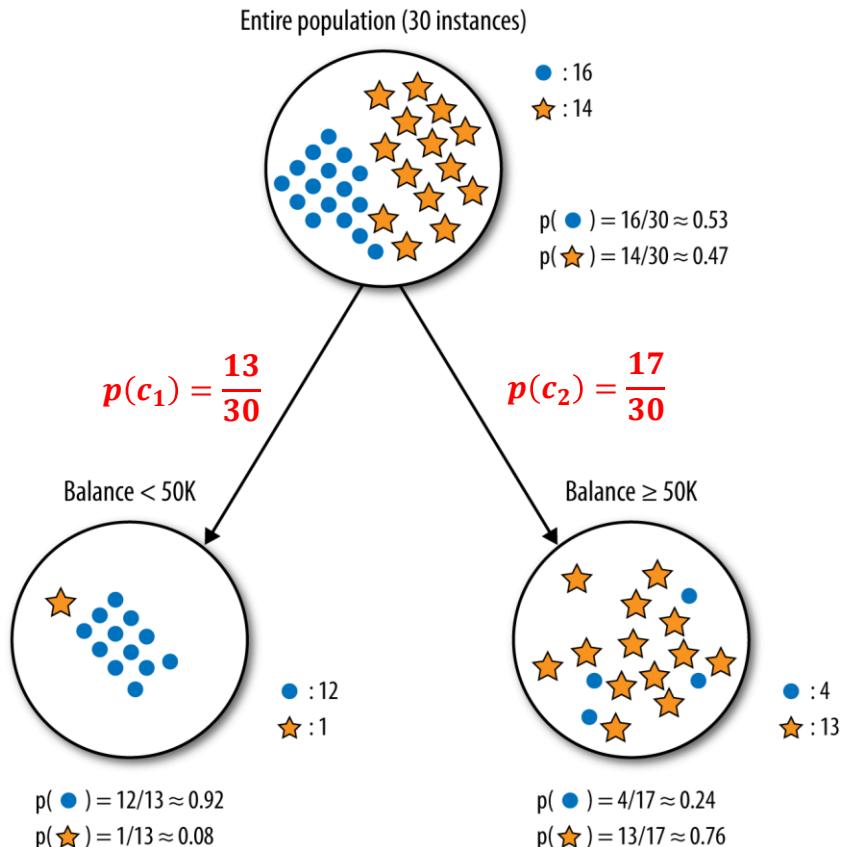
Entropy(parent)

- $= -\frac{16}{30} \times \log_2 \left(\frac{16}{30} \right) - \frac{14}{30} \times \log_2 \left(\frac{14}{30} \right)$
- $= 0.997$
- Very impure

Entropy(c_2)

- $= -\frac{4}{17} \times \log_2 \left(\frac{4}{17} \right) - \frac{13}{17} \times \log_2 \left(\frac{13}{17} \right)$
- $= 0.79$

Step 2 – Weighted Entropy of the Children



Entropy of each child

- ▶ $\text{Entropy}(c_1) = 0.39$
- ▶ $\text{Entropy}(c_2) = 0.79$

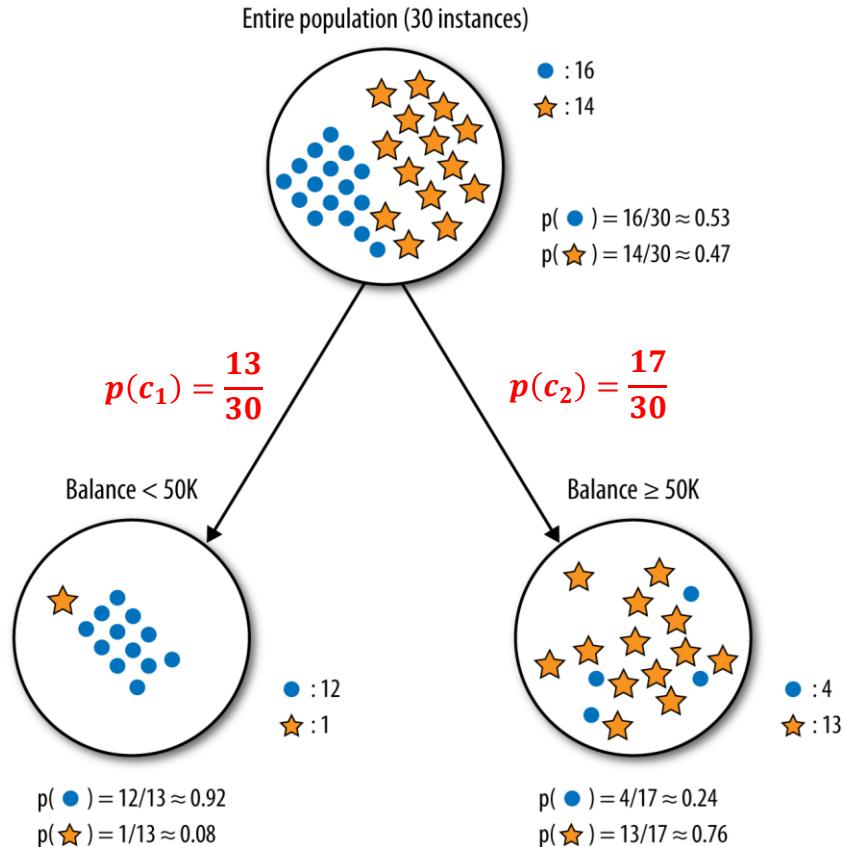
Proportion of instances in each child

- ▶ $p(c_1) = \frac{13}{30}$
- ▶ $p(c_2) = \frac{17}{30}$

Entropy (children)

- ▶ $= \frac{13}{30} \times 0.39 + \frac{17}{30} \times 0.79$
- ▶ $= 0.62$

Step 3 – Information Gain of *Balance*



Parent entropy = 0.99
Children entropy = 0.62

IG(Balance $\geq 50K$)

- ▶ $= 0.99 - 0.62$
- ▶ $= 0.37$

Compute Information Gain of *Residence*

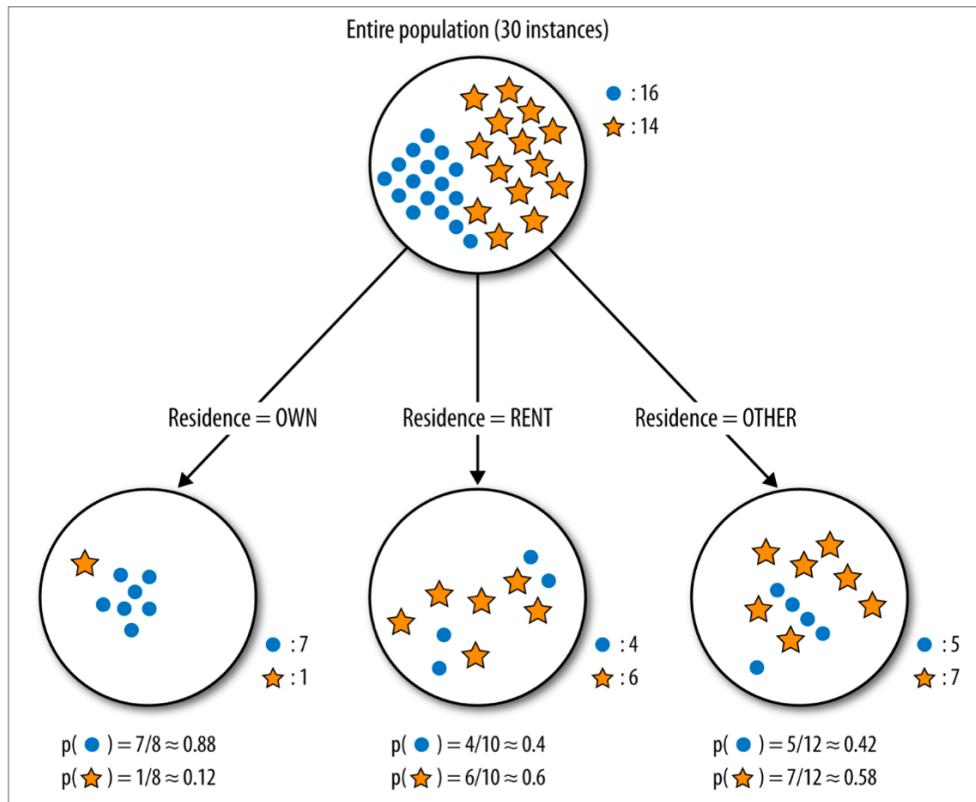


Figure 3-5. A classification tree split on the three-valued Residence attribute.

- *Entropy (parent)* ≈ 0.99
- *Entropy (Residence=OWN)* ≈ 0.54
- *Entropy (Residence=RENT)* ≈ 0.97
- *Entropy (Residence=OTHER)* ≈ 0.98
- *Entropy(children)* ≈ 0.86

...

$$IG(\text{Residence}) \approx 0.13$$

Conclusion:

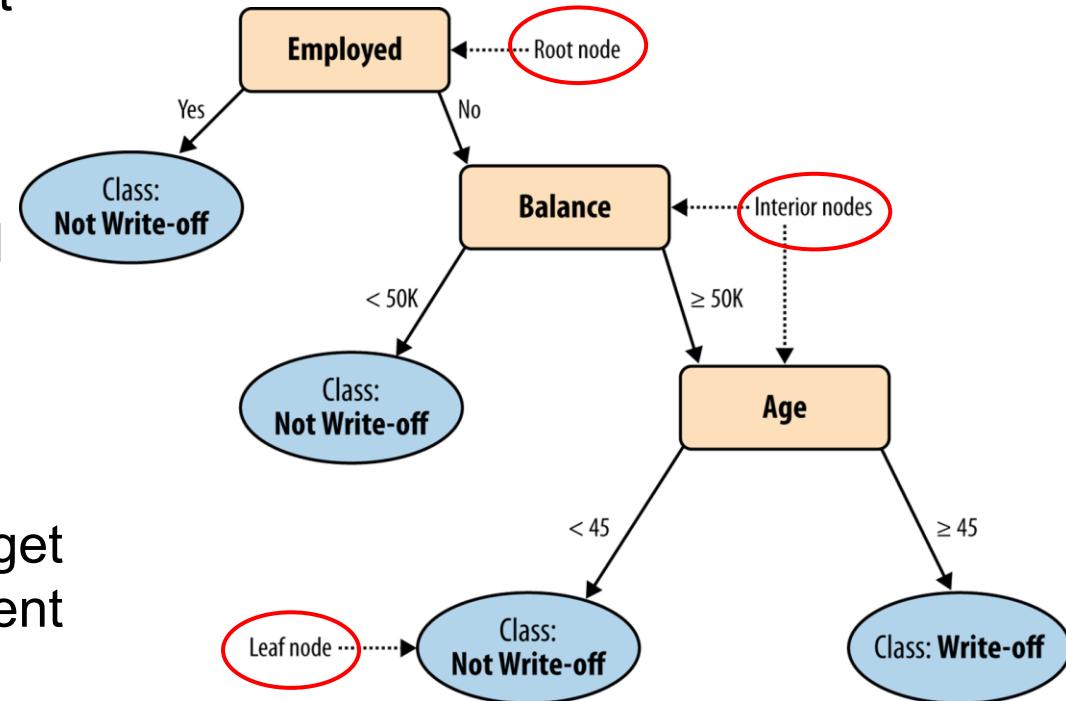
Balance is more informative than *Residence* in predicting loan write-off.

Train a Tree Model

- ▶ How do we build a classification tree?
 - ▶ **Divide-and-conquer** approach
 - ▶ Take each subgroup and **recursively** apply **feature/attribute selection** to find the best one for segmentation.
 - ▶ When shall we stop?
 - ▶ The leaf nodes are **pure**.
 - ▶ There is **no more features** for splitting.
 - ▶ **Avoid over-fitting** by controlling the size of the tree.

Tree Model: a Set of Rules

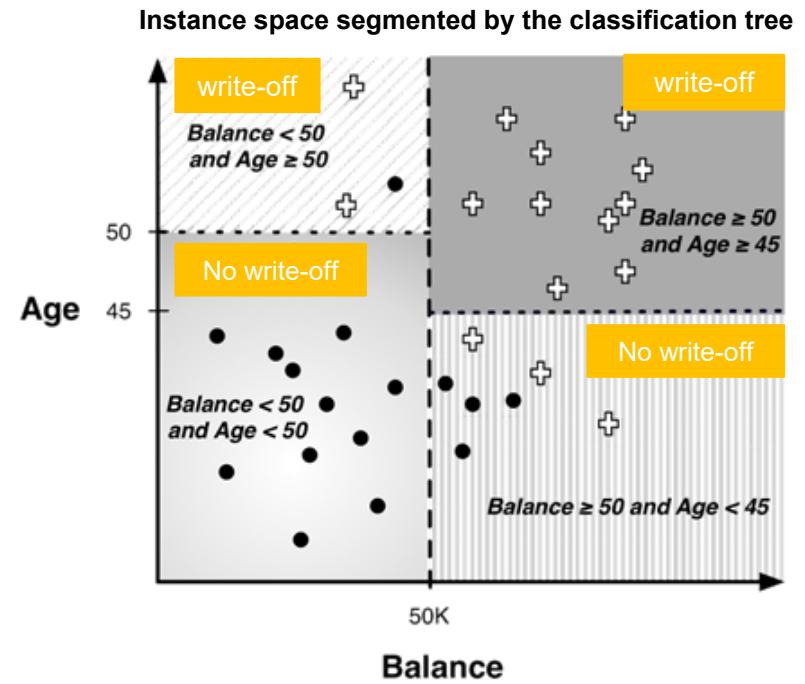
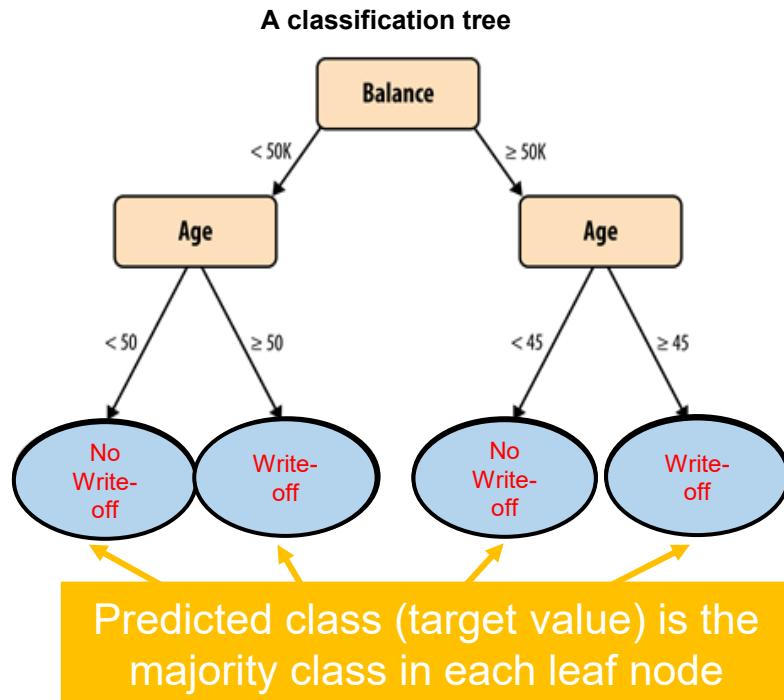
- ▶ **Root node:** the starting point
- ▶ **Interior node:** middle parts connecting nodes above and below
- ▶ **Terminal (leaf) node:** the classification outcome for target variable, each leaf is a segment
- ▶ **Branches** are defined by distinct values of the features



Unemployed clients with high outstanding balance and age above 45 tend to have their loans written off.

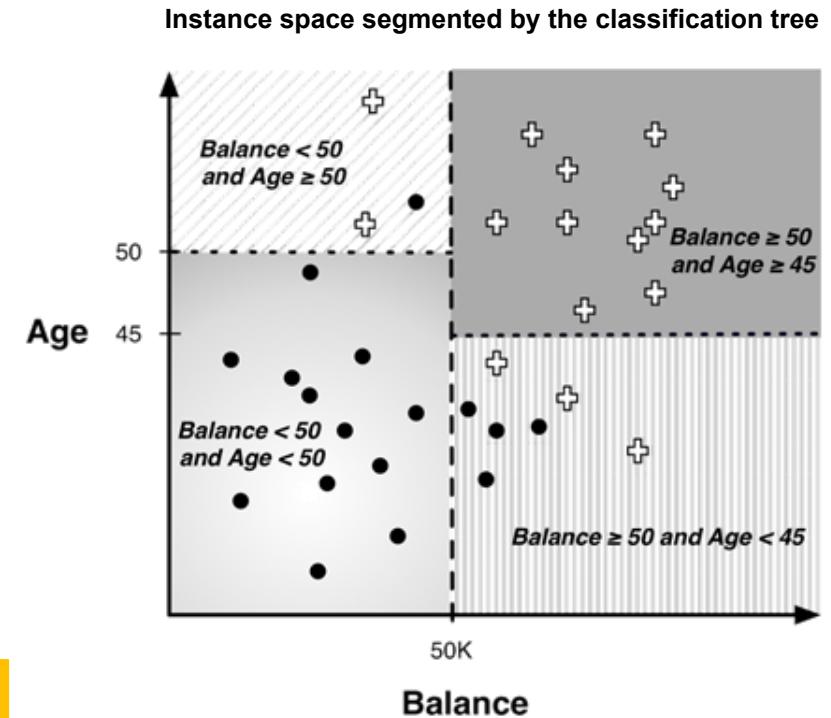
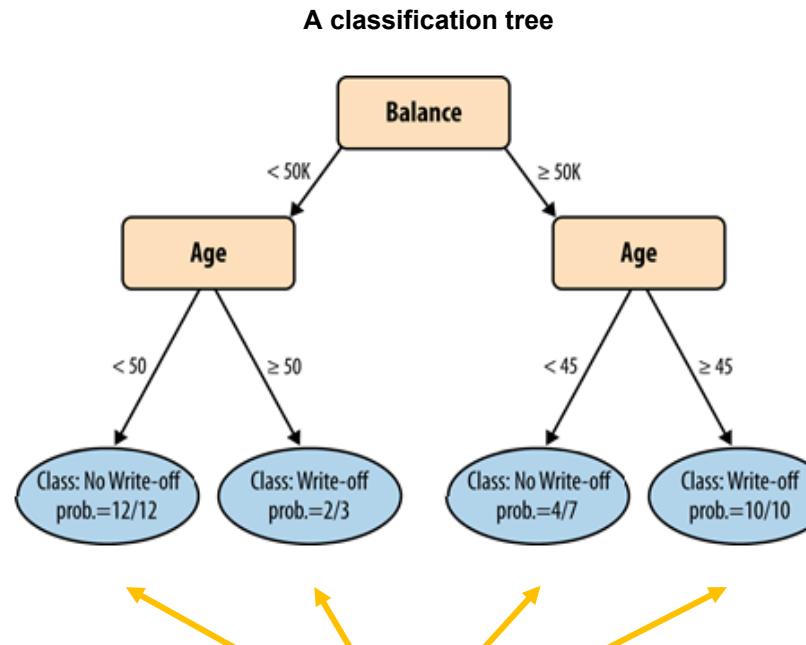
Class Prediction

- With a tree model trained, the segmentation result can be visualized in an **instance space**:
 - + write-off
 - no write-off
- Lines separating the regions are **decision boundaries**



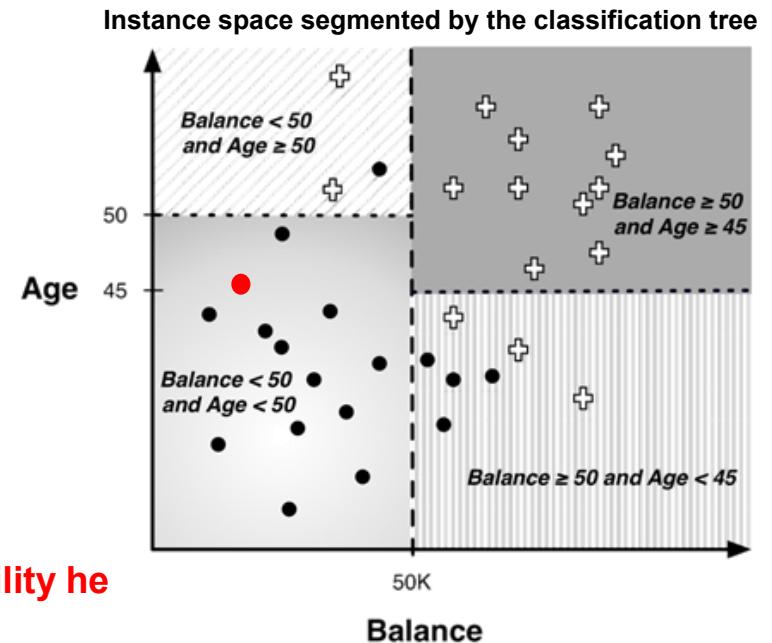
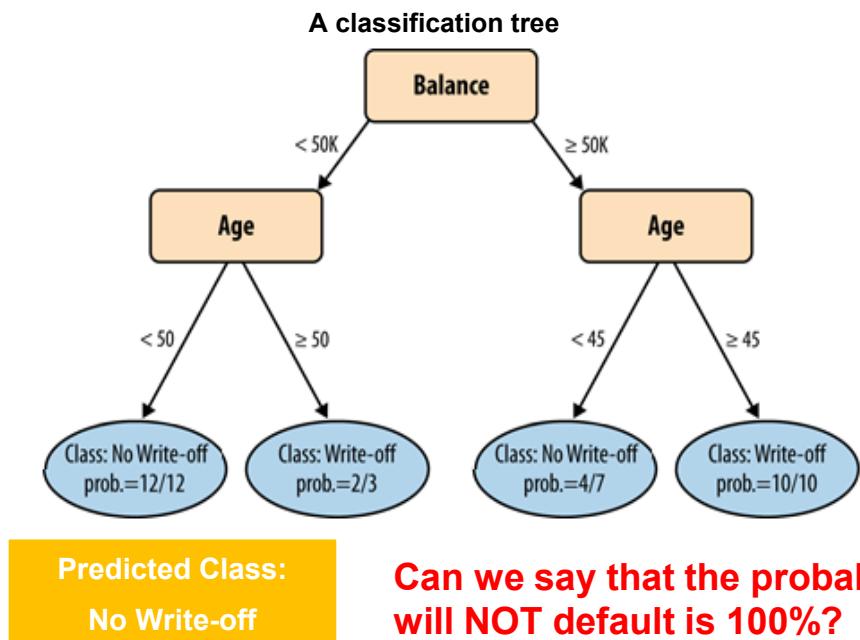
Class Probability Estimation

- ▶ How do we estimate the class probability?
- ▶ The segmenting result in some leaf nodes are impure.



Predict a New Customer's Loan Status

- ▶ For the loan write-off case, if a new customer:
 - ▶ Balance < 50K and age < 50
- ▶ What is the predicted class for this customer?
 - ▶ How likely will he default on the loan (thus be written off)?



Two Probability Estimation Methods

- ▶ Assume a leaf node with n positive, m negative instances:

Frequency-based estimate

- ▶ Probability that a new instance is positive is

$$\frac{n}{n+m}$$

Laplace correction

- ▶ Smoothed version of the frequency-based estimate

$$\frac{n+1}{n+m+2}$$

- ▶ Why a Laplace corrected probability is necessary?
 - ▶ With a bigger data size, we are more confident with the result.

(n, m)	$(2, 0)$	$(20, 0)$
Frequency	$\frac{2}{2} = 1$	$\frac{20}{20} = 1$
Laplace	$\frac{3}{4} = 0.75$	$\frac{21}{22} = 0.95$

Effect of Laplace Correction

- ▶ Consider two groups where 2/3 instances being positive:
 - ▶ A: $(n, m) = (4, 2)$, total 6
 - ▶ B: $(n, m) = (18, 9)$, total 27
- ▶ Frequency-based probability of positive class
 - ▶ Solid constant line = 2/3
- ▶ Laplace corrected probability
 - ▶ Dashed curve line
 - ▶ A: $p_{(c)} = \frac{4+1}{6+2} = 0.6250$
 - ▶ B: $p_{(c)} = \frac{18+1}{27+2} = 0.6552$

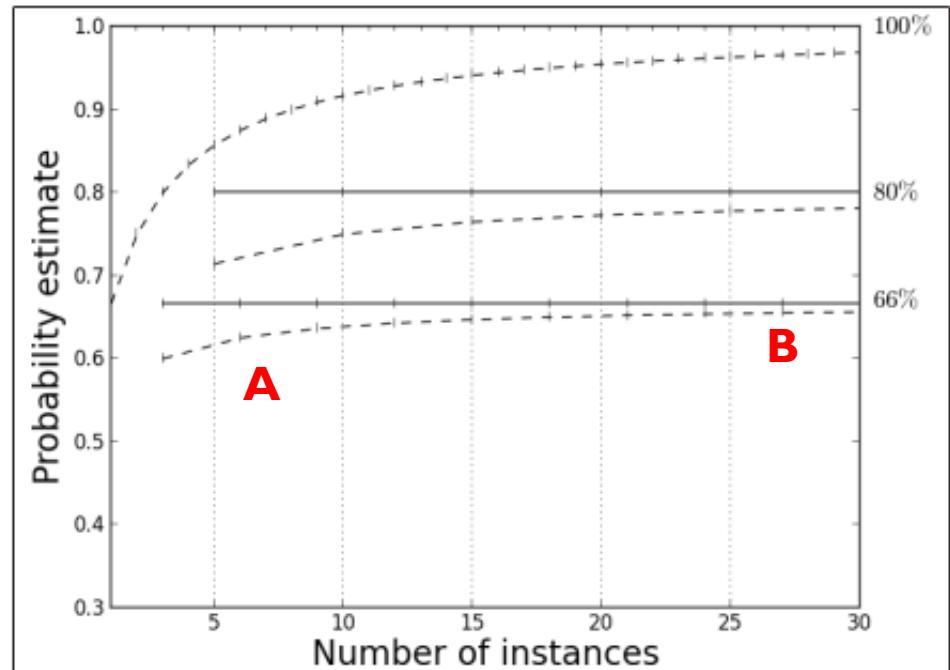


Figure 3-16. The effect of Laplace smoothing on probability estimation for several instance ratios.

As the number of total instances increases, Laplace-corrected probability converges to the frequency-based estimate.

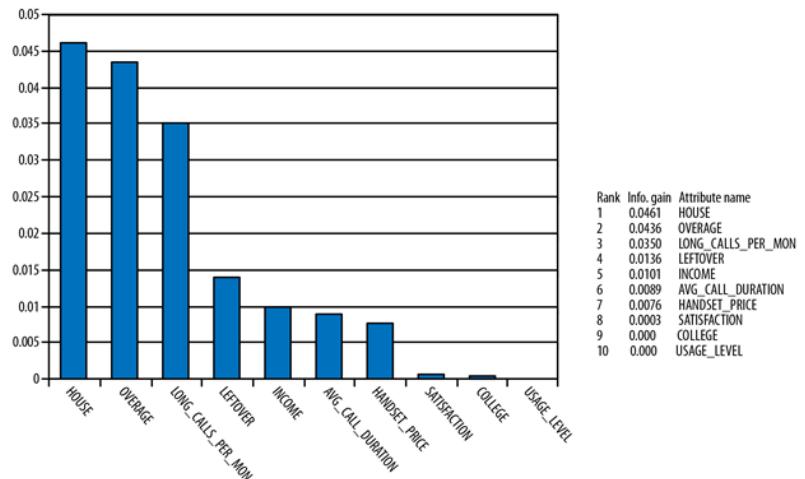
Compare Importance of Features

- ▶ Predicting **customer churn** based on their features.



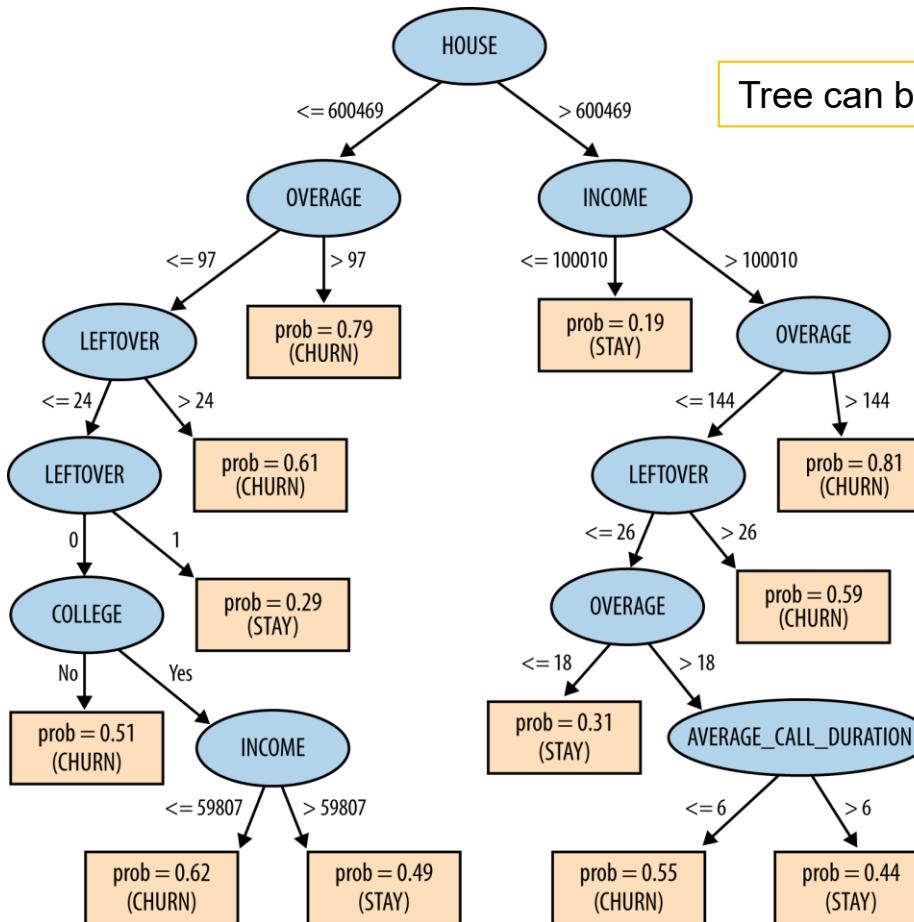
- ▶ Here is a list of features ranked by their **Information Gain**, evaluated on **entire dataset**.
 - ▶ Can we use these features one by one according the order?

Variable	Explanation
COLLEGE	Is the customer college educated?
INCOME	Annual income
OVERAGE	Average overcharges per month
LEFTOVER	Average number of leftover minutes per month
HOUSE	Estimated value of dwelling (from census tract)
HANDSET_PRICE	Cost of phone
LONG_CALLS_PER_MONTH	Average number of long calls (15 mins or over) per month
AVERAGE_CALL_DURATION	Average duration of a call
REPORTED_SATISFACTION	Reported level of satisfaction
REPORTED_USAGE_LEVEL	Self-reported usage level
LEAVE (<i>Target variable</i>)	Did the customer stay or leave (churn)?



Decision Tree

Same continuous features can be used more than once in the same branch.



Tree can be asymmetric.

Can categorical feature be used more than once in the same branch? No

There can be misclassification in the leaf nodes.

Leaf nodes from the same feature aims to have different class values.

Decision Trees

Pros:

- ▶ Easily visualized and interpreted.
- ▶ No feature normalization or scaling needed.
- ▶ Work well with datasets with a mixture of variable types (continuous, categorical).

Cons:

- ▶ Even after tuning, decision trees can often still **overfit**.
- ▶ Usually need an **ensemble of trees** for better generalization performance.