



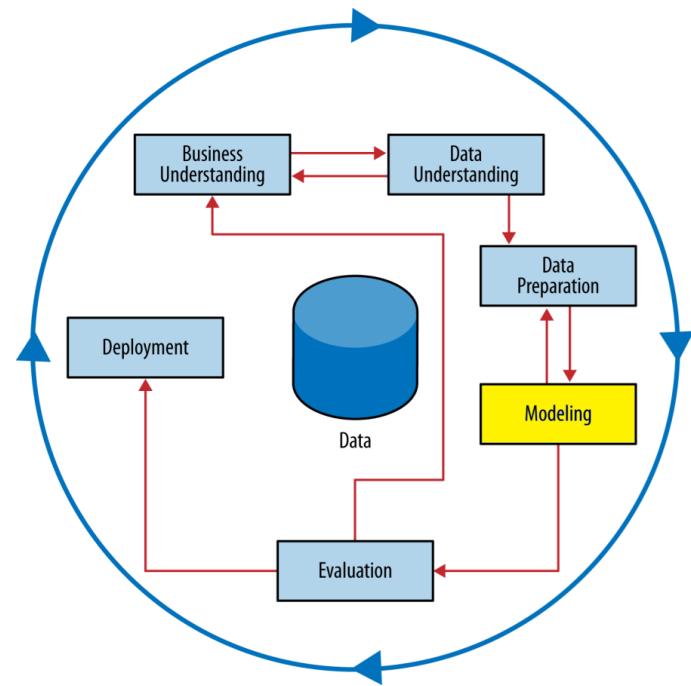
# **Fit a Model to Data: Logistic Regression and SVM**



**PF4**

# Learning Goals

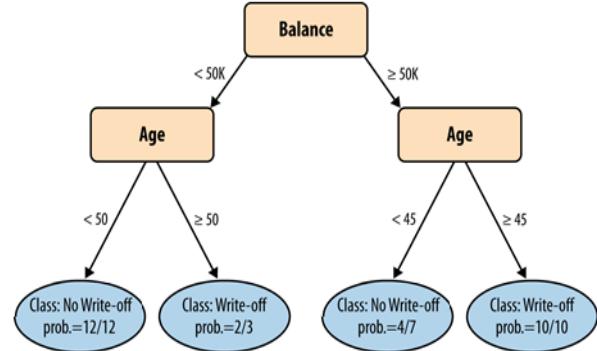
- ▶ Fit a model to the data
  - ▶ Learn the **optimal** parameters from the training data
- ▶ **Objective functions**
  - ▶ Is it the best fit with the data and the data mining goal?
- ▶ Supervised classification
  - ▶ Logistic regression
  - ▶ Support vector machine



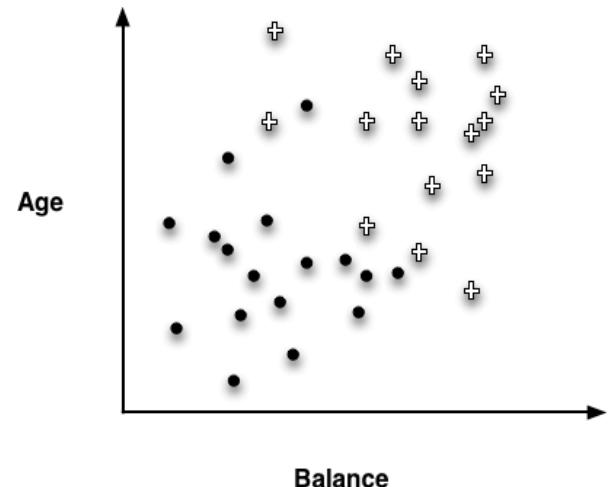
# Recap: Decision Tree

## ▶ Decision Tree

- ▶ Segment instances into sub-groups, so that members in the same group have similar target value.
- ▶ Each group is pure to some extent.
  - ▶ Degree of (im)purity measured by **entropy**



- ▶ The **decision boundary** in tree models is perpendicular to the axes.
  - ▶ Can we segment the customer data differently?



# Supervised Classification

---

## Non-parametric approach (e.g., Decision Tree)

- ▶ Decision boundaries are **perpendicular** to the axes
- ▶ Classify instances recursively using **divide-and-conquer** approach
- ▶ The model is a set of rules.

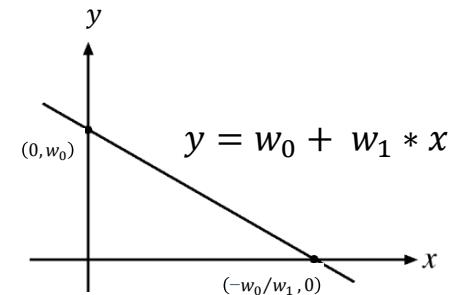
## Parametric approach (e.g., Logistic Regression, SVM)

- ▶ Linear classifiers can be in **any direction**
- ▶ Classify instances using **mathematical functions**
  - ▶  $f(x)$  is a weighted sum of various features  $x$

# Linear Classifier

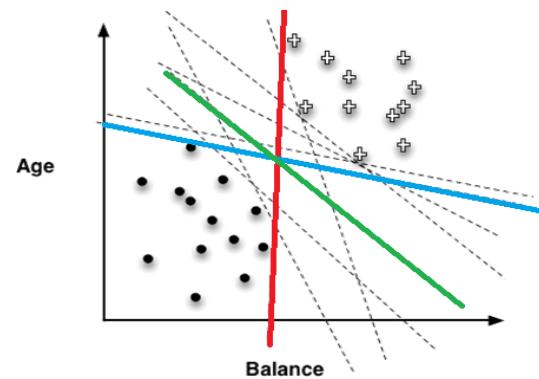
- ▶ A **Linear classifier** is weighted sum of various features  $x$ .
  - ▶ Weights refer to the coefficients.

- ▶ Recall the **mathematical function** for a simple linear regression?
  - ▶  $y = w_0 + w_1 x$ 
    - ▶  $w_0$  is the intercept (bias)
    - ▶  $w_1$  is the slope (coefficient)

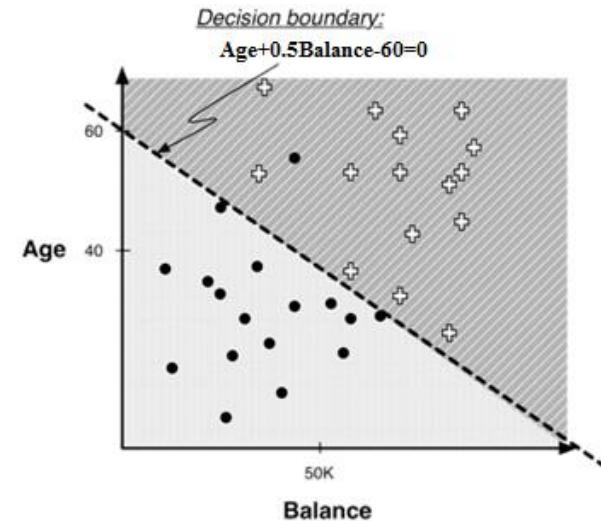
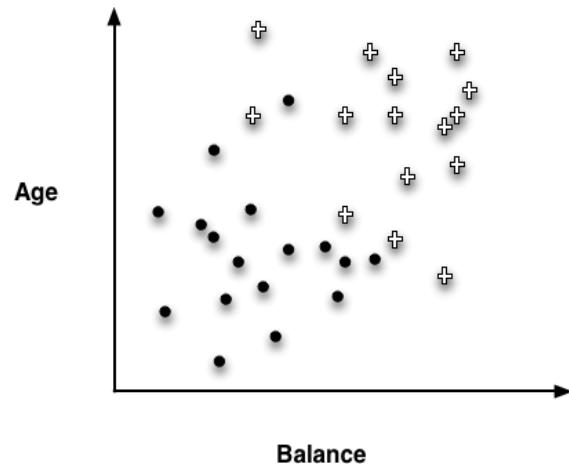


- ▶ We can also use a **straight line** to segment the customer data.
  - ▶ Each line is a linear mathematical function for *age* and *balance*.

$$Age = w_0 + w_1 \times Balance$$



# Linear Classifier



## ▶ Mathematical Function (Decision Boundary)

$$Age = -0.5 \times Balance + 60$$

$$\text{Age} + 0.5 \times \text{Balance} - 60 = 0$$

## ▶ Class Prediction

$$class = \begin{cases} \bullet & \text{if } Age + 0.5 \times Balance - 60 \leq 0 \\ + & \text{if } Age + 0.5 \times Balance - 60 > 0 \end{cases}$$

# Linear Classifier

$$f(X) = 2X_1 + 3X_2 + 1$$

- ▶ The decision boundary:
  - ▶  $f(X) = 2X_1 + 3X_2 + 1 = 0$
  - ▶ or  $X_2 = -\frac{2}{3}X_1 - \frac{1}{3}$
- ▶ To have  $f(X) = 0$ :
  - ▶ When  $X_1 = -\frac{1}{2}$ ,  $X_2 = 0$
  - ▶ When  $X_1 = 0$ ,  $X_2 = -\frac{1}{3}$

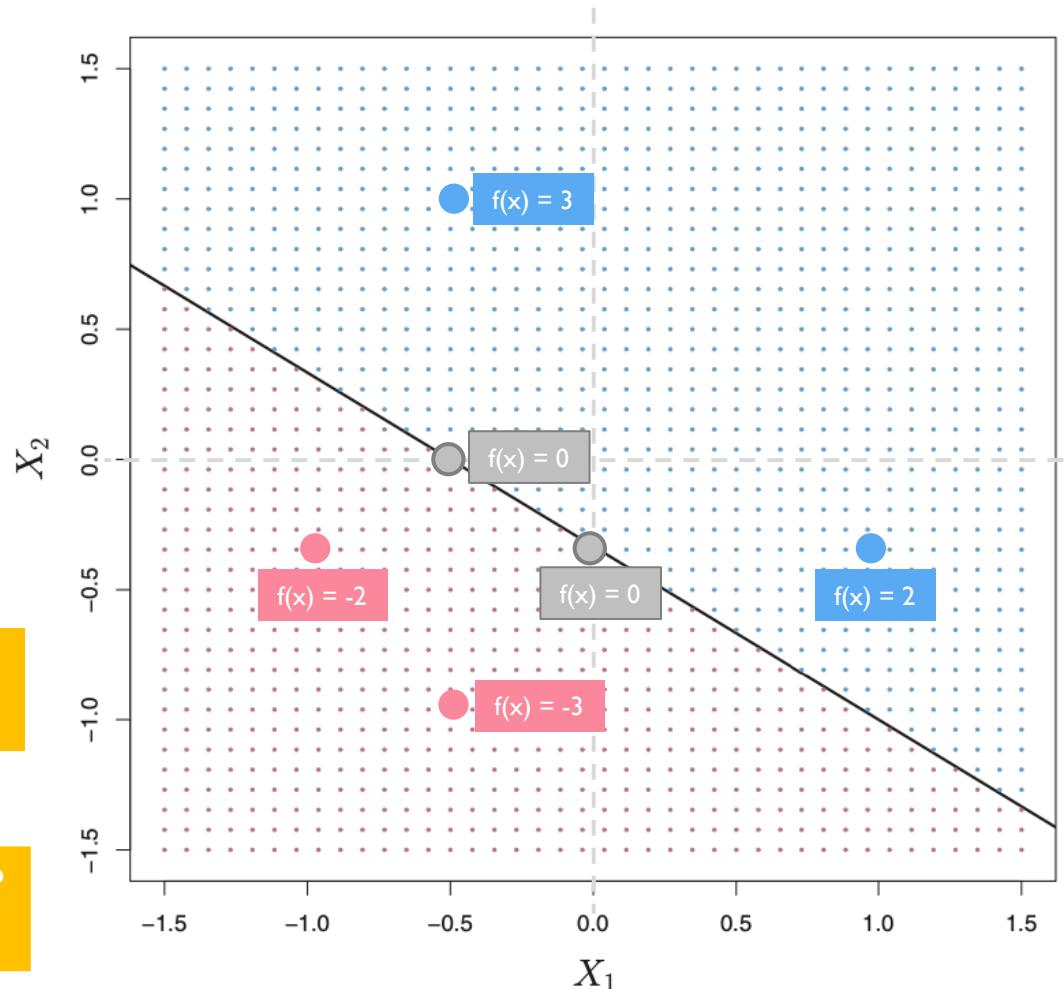
Calculate  $f(X)$  value if  $X_2 = 1$  or  $-1$ ?

No change for  $X_1$ .

▶ When  $X_1 = 0$ ,  $X_2 = -\frac{1}{3}$

Calculate  $f(X)$  value if  $X_1 = -1$  or  $1$ ?

No change for  $X_2$ .

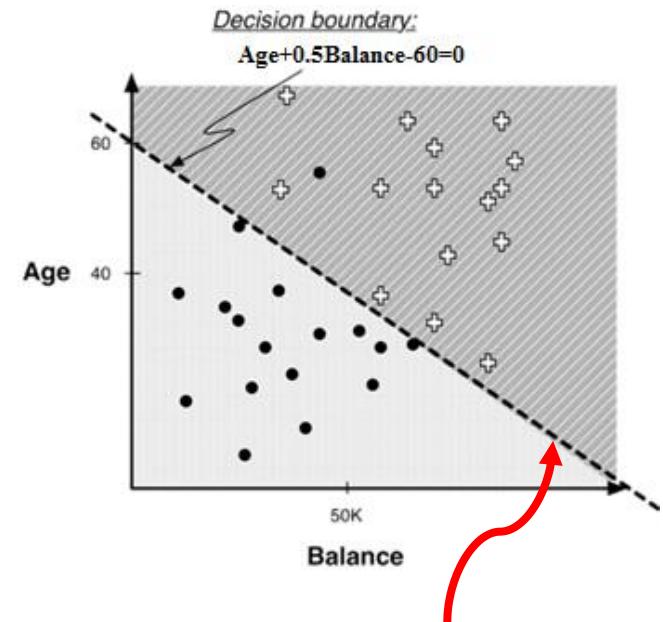


# Linear Classifier

- ▶ A general linear model

$$f(x) = w_0 + w_1x_1 + \cdots + w_px_p$$

- ▶  $|f(x)|$  measures the **distance** between an **instance** with features  $x$  and the **decision boundary**.
- ▶  $|f(x)|$  is NOT the perpendicular distance to decision boundary, but proportional to it.
- ▶  $|f(x)|$  reflects the degree of **certainty** that an instance belongs to a class.
  - ▶ Example: loan write-off
    - + write-off
    - not write-off
  - ▶  $|f(x)|$  increases, certainty increases

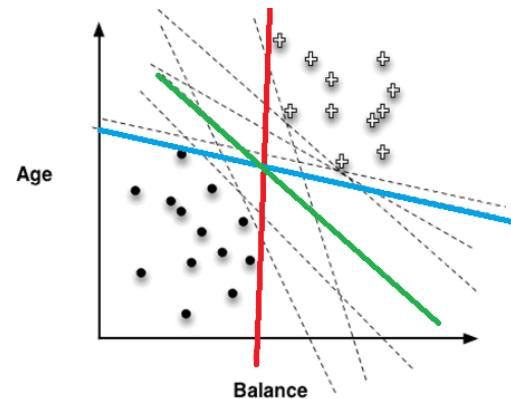


$f(x) = 0$ : **most uncertain**

Perpendicular distance of an instance to decision boundary:  $D = |f(x)| / \sqrt{\sum_{j=1}^p w_j^2}$

# Objective Function: Find the Best Line

- ▶ The optimal parameter  $w_p$  depends on the **objective function**.
  - ▶ An objective function represents what we want to achieve in a data mining task.
  - ▶ It usually tries to reduce the **error/loss** for all training instances.
    - ▶ There are various error/loss functions to choose.
- ▶ Like linear regression, both logistic regression and SVM optimize the parameters by **fitting a (linear) model to data**, the difference is in their objective functions.
  - ▶ recap: linear regression minimizes **sum of squared errors** (i.e., RSS).



$$\min \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

What is our goal in a classification task then?

# Logistic Regression: Probability Estimate

---

- ▶ Logistic regression applies a **linear mathematical function** to estimate the **class probability**.
- ▶ Binary classification:  $f(x)$  is used to estimate  $P(y = 1 | x)$ .

$$f(x) = w_0 + w_1x_1 + \cdots + w_px_p$$

- ▶ What is the range of the linear output  $f(x)$  ?  
[-∞, ∞]
- ▶ What is the range of a class probability  $P(y = 1 | x)$  ?  
[0, 1]
- ▶ How can we convert  $f(x)$  into  $P(y = 1 | x)$ ?

# Step 1: Odds and Probability

- ▶ The **odds (ratio)** of an event is the ratio of the probability of the event *occurring* vs. the probability of the event **not occurring**.
- ▶ Assume  $p$  is the probability of instance  $x_i$  belonging to the positive class, the **odds** of instance  $x_i$  being positive is

$$\frac{p}{1-p} = \frac{P(y=1 | x)}{P(y=0 | x)}$$

Odds is ranged from  $[0, \infty]$

Probability ( $p$ )	Odds ( $\frac{p}{1-p}$ )
0.5	50:50 (or 1)
0.9	90:10 (or 9)
0.999	999:1 (or 999)
0.01	1:99 (or 0.0101)
0.001	1:999 (or 0.001001)

## Step 2: Odds and Log-odds

- ▶ The natural logarithm of odds is ranged from  $[-\infty, \infty]$ .
- ▶ Natural logarithm takes base  $e$ , a mathematical constant 2.71...

Probability ( $p$ )	Odds ( $\frac{p}{1-p}$ )	Log-odds ( $\ln(\frac{p}{1-p})$ )
0.5	50:50 (or 1)	0
0.9	90:10 (or 9)	2.19
0.999	999:1 (or 999)	6.9
0.01	1:99 (or 0.0101)	-4.6
0.001	1:999 (or 0.001001)	-6.9

- ▶ The linear value  $f(x)$  estimates the log-odds that instance  $x$  belong to the positive class:

$$f(x) = w_0 + w_1 x_1 + \dots + w_p x_p = \ln \left( \frac{p}{1-p} \right)$$

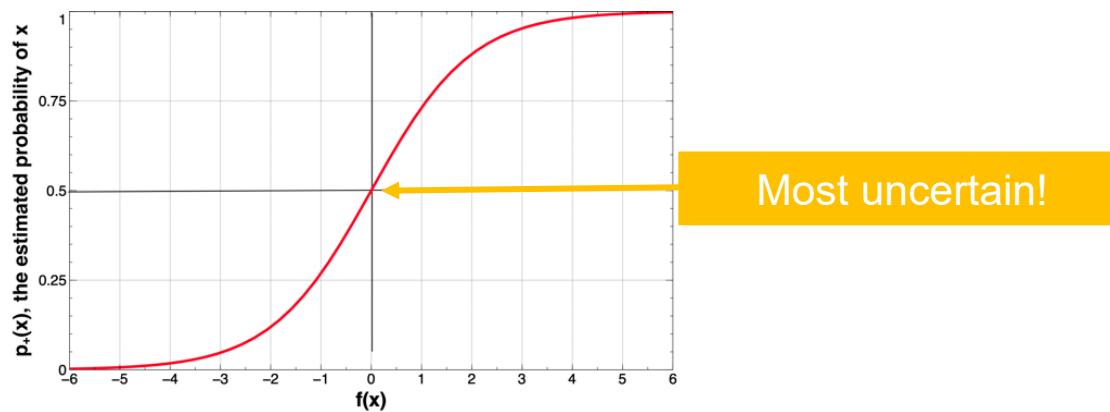
Both  $\ln(x)$  or  $\log(x)$  represent natural logarithm of  $x$

# Logistic Regression: Probability Estimate

- ▶ **Logistic Function**, also named sigmoid function, transforms log-odds to class probability.
- ▶ Class probability is a logistic function of log-odds (i.e.,  $f(x)$ ).

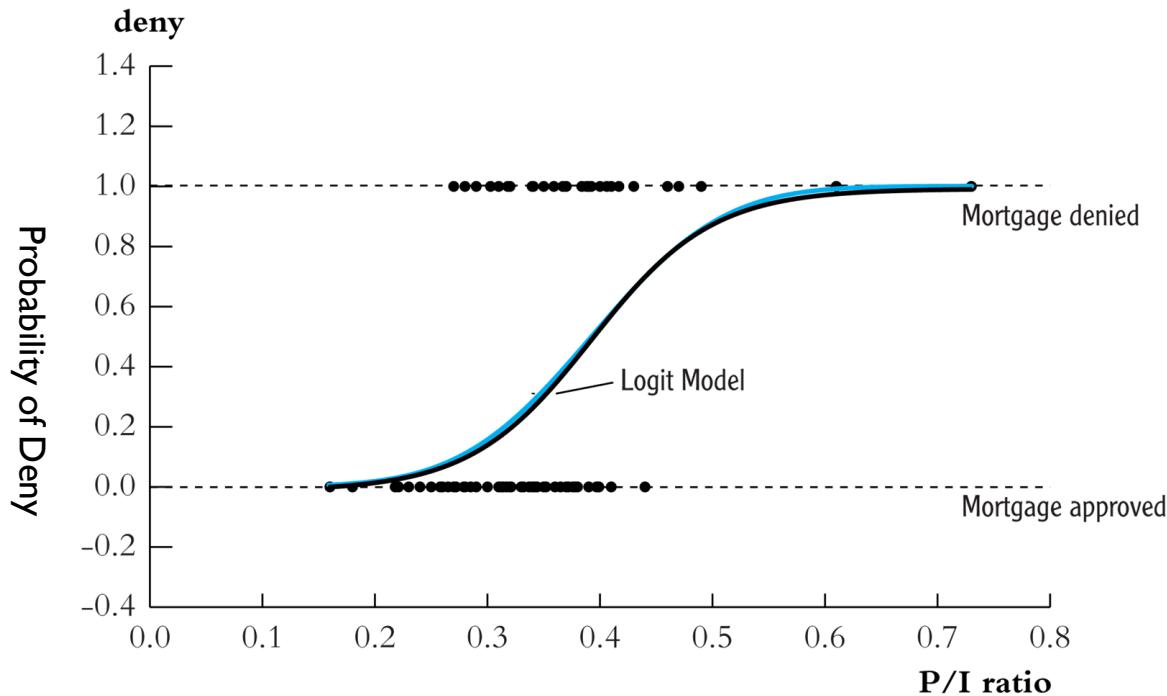
$$P(y = 1|f(x)) = \frac{e^{f(x)}}{1+e^{f(x)}} \text{ or } \frac{1}{1+e^{-f(x)}}$$

- ▶ When  $f(x) = 0$ , what is  $P(y = 1)$ ?
- ▶ How shall we interpret it?



# Example: Mortgage Application

- ▶ Bank's decision: deny a mortgage application?
  - ▶ **Target:** Deny = reject (1) or approve (0) the application.
  - ▶ **Feature:** *P/I ratio* measures the ratio of debt to income for a person.



# Example: Mortgage Application

---

- ▶ The logistic regression model estimates  $w_0 = -1$  and  $w_1 = 2.5$

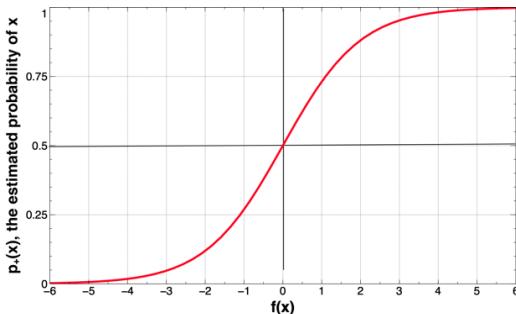
$$f(x) = -1 + 2.5 \times PI\ ratio$$

$$P(deny = 1|PI\ ratio) = \frac{1}{1 + e^{-f(x)}}$$

- ▶ Interpret the coefficient ( $w_1$ ): 1 unit increase in PI ratio is associated with 2.5 unit increase in log-odds of loan rejection.
- ▶ When P/I ratio = 0.3:
  - ▶ The log-odds  $f(x) = -1 + 2.5 \times 0.3 = -0.25$
  - ▶ The probability  $P(deny = 1|PI\ ratio) = \frac{1}{1+e^{0.25}} \approx 0.44$ 
    - ▶ There is 44% chance that the application will be rejected.

# Objective Function: Logistic Regression

- ▶ What is our goal in a binary classification task?
  - ▶ For positive instances (i.e., actual class  $y_i = 1$ ), we expect
    - ▶  $\hat{y}_i = p(+)$  is as close to 1 (i.e.,  $y_i$ ) as possible.
  - ▶ For negative instances (i.e., actual class  $y_i = 0$ ), we expect
    - ▶  $\hat{y}_i = p(+)$  is as close to 0 (i.e.,  $y_i$ ) as possible. This means,  $p(-)$  or  $1 - \hat{y}_i$  is as close to 100% as possible.



**Maximum Likelihood Estimation:**  
maximize the estimated probability of each instance's actual class.

- ▶ **Objective Function:**

$$\min_w \sum_{i=1}^n -y_i \ln(\hat{y}_i) - (1 - y_i) \ln(1 - \hat{y}_i)$$

Log-loss or Binary Cross-entropy

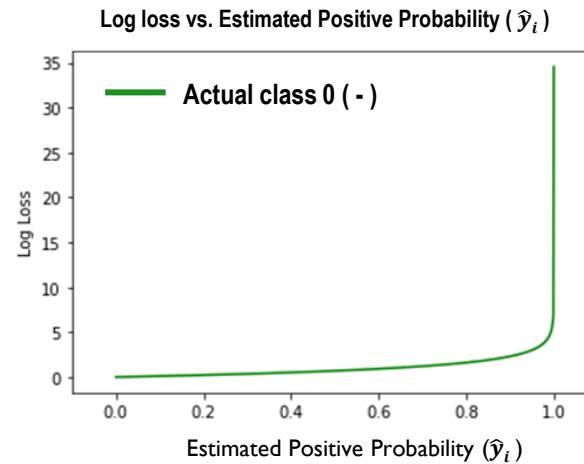
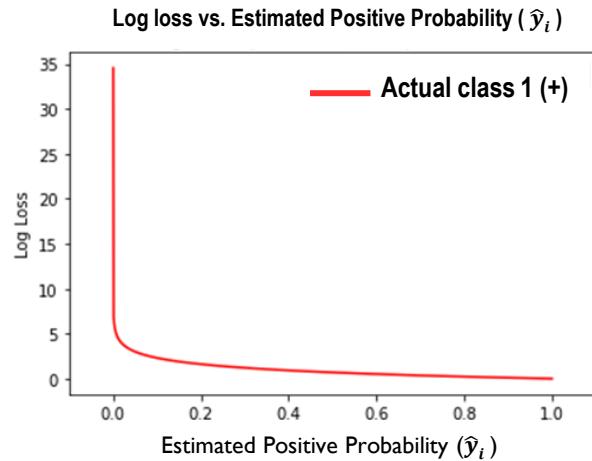
$$\text{Log-loss} = -y_i \ln(\hat{y}_i) - (1 - y_i) \ln(1 - \hat{y}_i)$$

# Objective Function: Logistic Regression

- For any instance, if  $p(+)$  (i.e.,  $\hat{y}_i$ ) goes farther away from its actual class label (i.e.,  $y_i$ ), its **log loss** is greater.
  - For a positive instance (i.e.,  $y_i = 1$ ): the farther away  $\hat{y}_i$  is from 1 (actual class  $y_i$ ), the greater its log loss is.
  - For a negative instance (i.e.,  $y_i = 0$ ): the farther away  $\hat{y}_i$  is from 0 (actual class  $y_i$ ), the greater its log loss is.

$$\text{Log-loss} = -\ln(\hat{y}_i)$$

$$\text{Log-loss} = -\ln(1-\hat{y}_i)$$



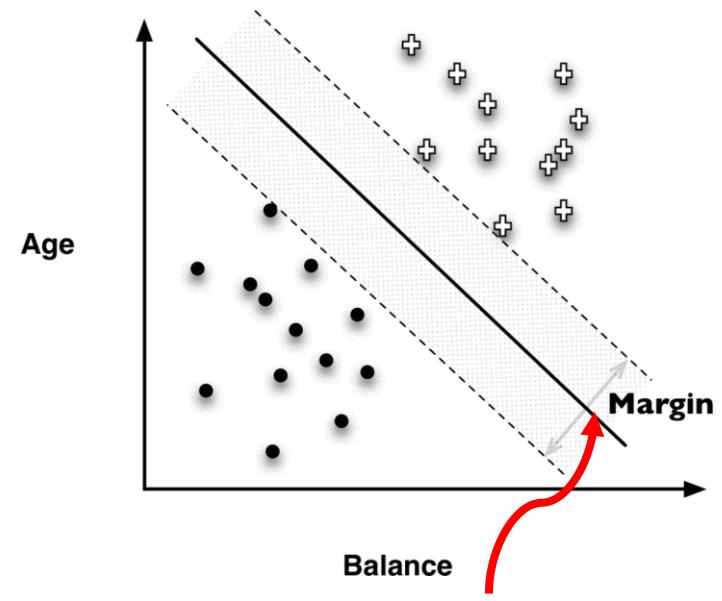
- By **minimizing the sum of log loss** for all training instances, logistic regression maximizes the **estimated probability of actual class** for all.

# Support Vector Machine

- ▶ **Support vector machines** are also linear classifiers as well.

$$f(x) = w_0 + w_1x_1 + \dots + w_px_p = 0$$

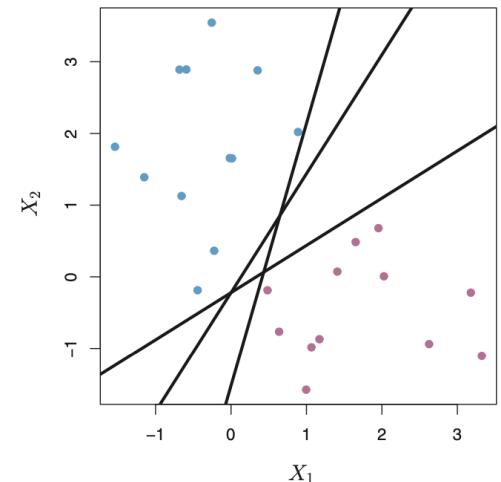
- ▶ Instead of separating data with a line directly, SVM also fit the **widest bar** (i.e., margin) between the two classes.
  - ▶ Why the bar is necessary?



$f(x) = 0$ : most uncertain

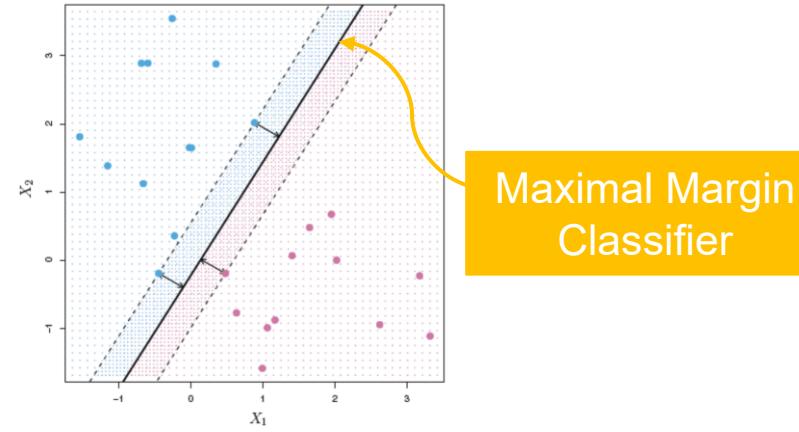
# Why a Margin is Necessary?

- ▶ Given a 2D feature space  $(X_1, X_2)$ , find a **separating hyperplane** (i.e., **decision boundary**) to separate two classes  $y_i = \{\text{Yes}, \text{No}\}$ .
  - ▶ The decision boundary is  $f(x) = w_0 + w_1X_1 + w_2X_2 = 0$ 
    - ▶  $f(x) > 0$ , predict Yes
    - ▶  $f(x) < 0$ , predict No
- ▶ However, there can be multiple separating hyperplanes for a dataset.
  - ▶ in the right graph, all lines can be shifted a tiny bit, without touching any instance.
- ▶ With training instances at least **certain distance away** (i.e., margin) from the hyperplane, we can avoid the **uncertain area**.
  - ▶ The margin prevents the hyperplane from being overly influenced by noises and therefore **avoid overfitting**.



# Road to SVM: Maximal Margin Classifier

- ▶ **Hard margin**: no training instance shall violate the margin lines.
  - ▶ All instances are separable – no misclassification.
  - ▶ Instances are at least certain distance away from their margin line.
- ▶ To find the best hyperplane (i.e., the model):
  - ▶ For each hyperplane, find the **margin width**: i.e., the shortest perpendicular distance between instances and the hyperplane.
    - ▶ Need to compare the perpendicular distance for all instances (to hyperplane).
  - ▶ The best hyperplane is the one with the **maximized margin**.
- ▶ **Support Vectors**: instances on the margin lines (i.e., dashed lines).
  - ▶ **Equidistant** to the hyperplane.
  - ▶ Only movement of support vectors affects the model.
    - ▶ Movement of other instances does not.



# Objective Function: Maximal Margin Classifier

$$\max_{w_1 \dots w_p, M} M \quad \text{subject to} \quad \sum_{j=1}^p w_j^2 = 1$$

Here  $y_i$  takes either 1 or -1, and  $\hat{y}_i = f(x)$

- For positive instances ( $y_i = +1$ ):  $\hat{y}_i = w_0 + w_1 x_{1i} + \dots + w_p x_{pi} > 0$
- For negative instances ( $y_i = -1$ ):  $\hat{y}_i = w_0 + w_1 x_{1i} + \dots + w_p x_{pi} < 0$

For separable training instances  $x_i$ :

$$y_i \hat{y}_i \geq M \quad \text{Margin Width (always } > 0\text{)}$$

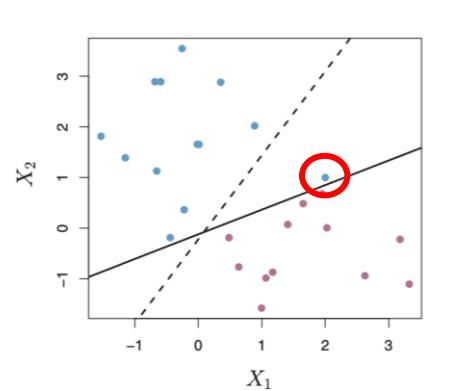
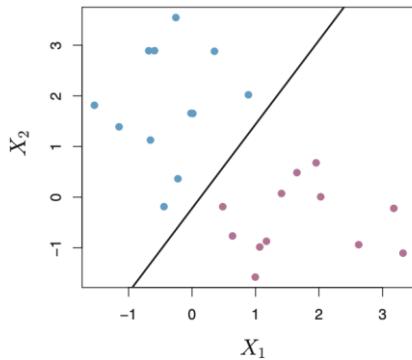
Note: Maximal Margin Classifier assumes all instances are separable.

$y_i \hat{y}_i = |\hat{y}_i|$  : the *perpendicular distance* between instance  $x_i$  to the hyperplane, given  $\sum_{j=1}^p w_j^2 = 1$

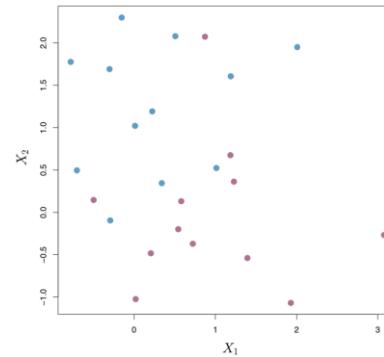
Perpendicular distance of  $x_i$  to the hyperplane:  $D_i = |\hat{y}_i| / \sqrt{\sum_{j=1}^p w_j^2}$

# Drawbacks of Maximal Margin Classifier

- Maximal Margin Classifier is **extremely sensitive** to small changes in training data and tend to **overfit**.
- According to its objective function, each instance should not only be on the correct side of **the hyperplane**, but also the correct side of **its margin line**.

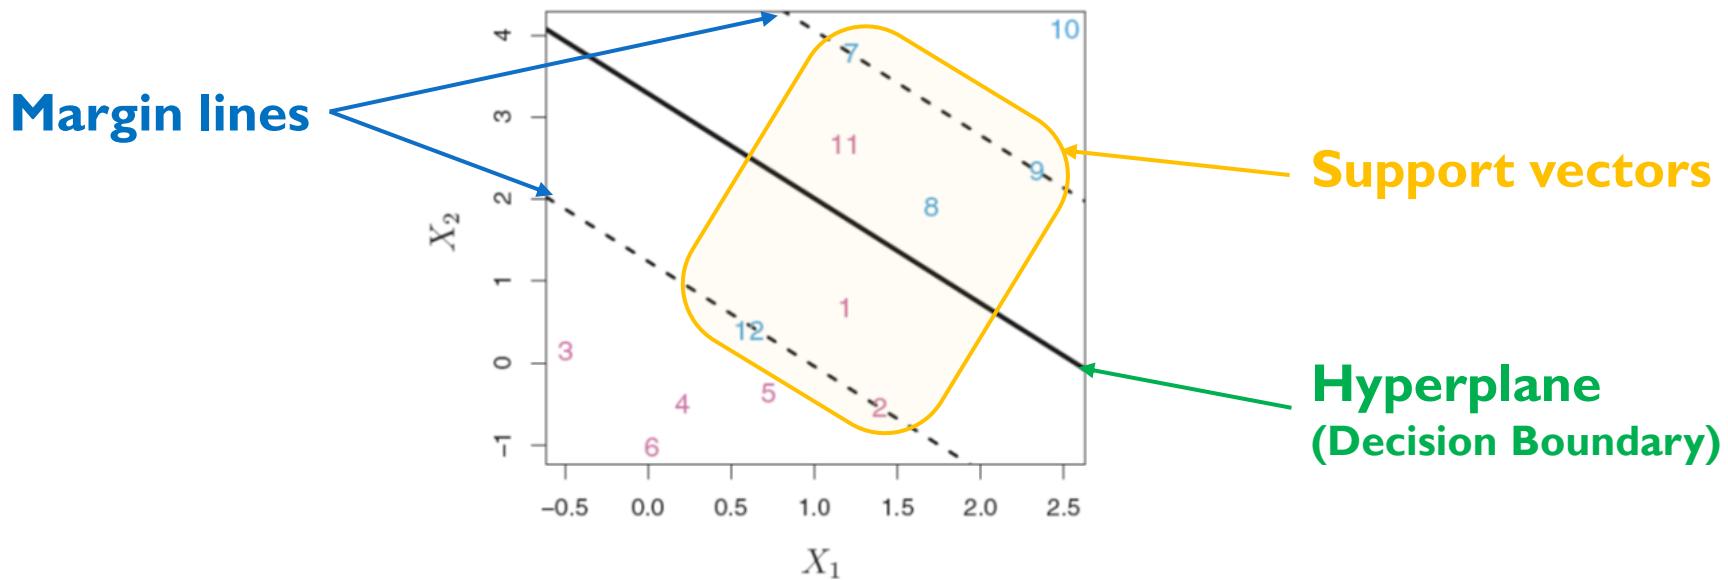


- The separating hyperplane **may not exist**.
  - When training instances are non-separable, there is **NO** maximal margin classifier with margin width  $> 0$ .



# Support Vector Classifier

- ▶ **Support vector classifier** relies on **soft margin**, which allows some training instances to be the incorrect side of its margin line.
  - ▶ **Support vectors** are training instances violating its margin line.
    - ▶ Support vectors bring errors to objective function and affect the hyperplane.
  - ▶ Only instances violating the hyperplane will be misclassified.
    - ▶ Instances that violating its margin line but not the hyperplane bring errors to objective function but will NOT be misclassified.



# Objective Function: SVC

$$\max_{w_1 \dots w_p, \xi_i, M} M \quad \text{subject to} \quad \sum_{j=1}^p w_j^2 = 1$$

Here  $y_i$  takes either 1 or -1, and  $\hat{y}_i = f(x)$

- $\hat{y}_i = w_0 + w_1 x_{1i} + \dots + w_p x_{pi}$  can be any value

$$y_i \hat{y}_i \geq M(1 - \xi_i) \quad \text{where } \xi_i \geq 0$$

**Hinge loss** allows instance  $x_i$  to be at distance  $\xi_i$  from its margin line.

- $\xi_i = 0$ : in the correct side of its margin line.
- $0 < \xi_i < 1$ : violate its margin line only (but not hyperplane)
- $\xi_i > 1$ : violate the hyperplane (**misclassification**)

► Rewritten as:

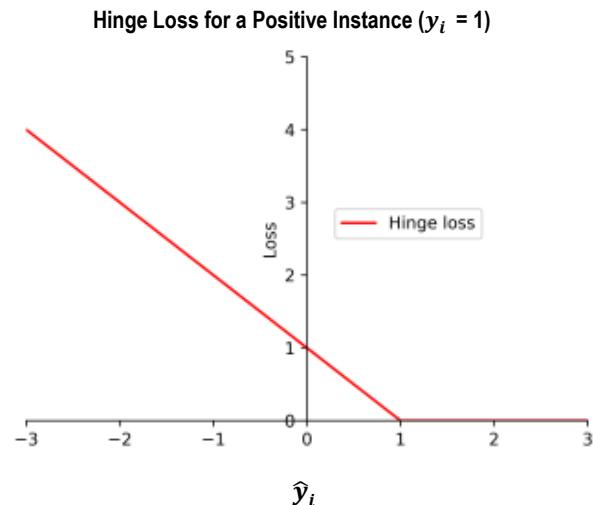
$$\min_w \sum_{i=1}^n \xi_i \quad \text{where } \xi_i = \max_i [0, 1 - y_i \hat{y}_i]$$

If  $\sum_{j=1}^p w_j^2 = 1$ , margin width  $M = 1 / \sqrt{\sum_{j=1}^p w_j^2} = 1$

# Hinge Loss

- ▶ Hinge loss ( $\xi$ ) is the error for instances violating its own margin line.
- ▶ Hinge loss is proportional to the distance between the instance and its margin line.

$$\xi_i = \max_i [0, 1 - y_i \hat{y}_i]$$



- ▶ For a positive instance ( $y_i = 1$ ):
  - ▶  $\hat{y}_i > 1$ : in the positive side and outside of positive margin - no loss.
    - ▶  $1 - y_i \hat{y}_i < 0 \rightarrow \xi_i = 0$
  - ▶  $0 < \hat{y}_i < 1$ : still in the positive side, violates positive margin only - loss occurs.
    - ▶  $0 < 1 - y_i \hat{y}_i < 1 \rightarrow \xi_i = 1 - y_i \hat{y}_i$  (ranged between [0,1])
  - ▶  $\hat{y}_i < 0$ : in the negative side, misclassification - greater loss.
    - ▶  $1 - y_i \hat{y}_i > 1 \rightarrow \xi_i = 1 - y_i \hat{y}_i$  (greater than 1)

Hinge loss for a negative instance ( $y_i = -1$ ) is symmetrical to that of a positive instance.

# Support Vector Machine

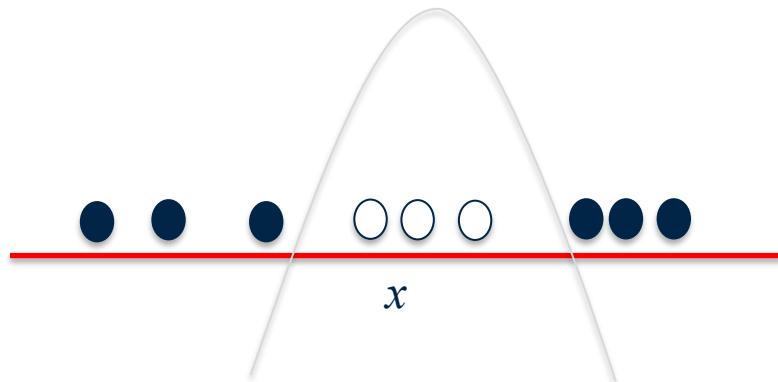
---

- ▶ **Support vector machine** is an extension of support vector classifier. It enlarges the original feature space by mapping them to higher dimensions, using *kernels*.
  
- ▶ **Linear kernel** → Support vector classifier
  - ▶ No exponential or interaction terms.
  
- ▶ **Polynomial kernel with degree**
  - ▶ e.g., a polynomial kernel maps the original 2D feature space  $(X_1, X_2)$  to a new feature space  $(X_1, X_2, X_1^2, X_2^2, X_1X_2)$ .
  
- ▶ **Other kernels**
  - ▶ Options includes “*rbf*”, “*sigmoid*”, “*precomputed*”.

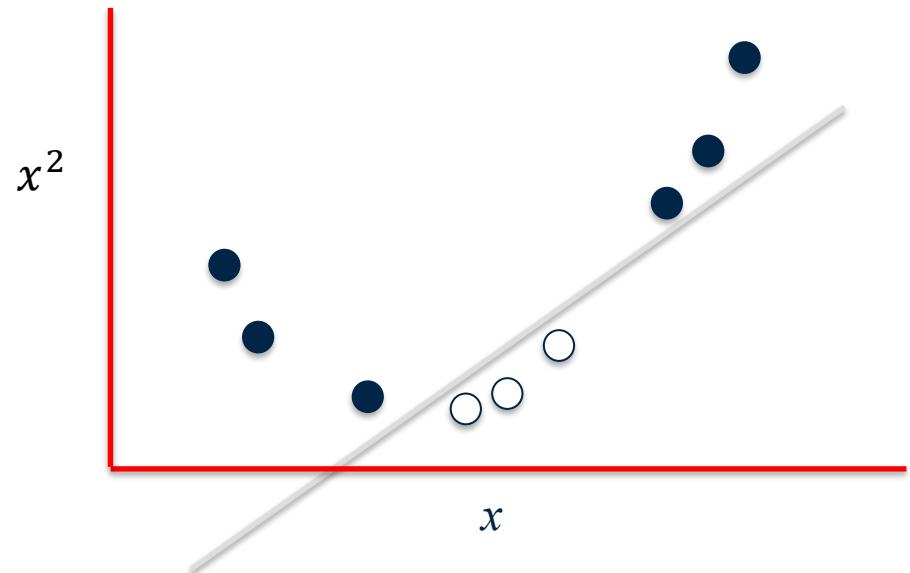
# Accommodating Nonlinearity

- Inseparable instances in 1D space (left) become linearly separable in a 2D feature space (right).
  - How would the hyperplane in 2D feature space look like in the original 1D feature space?

Original feature space ( $x$ )

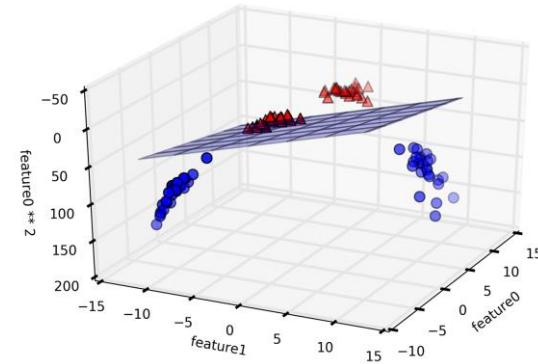
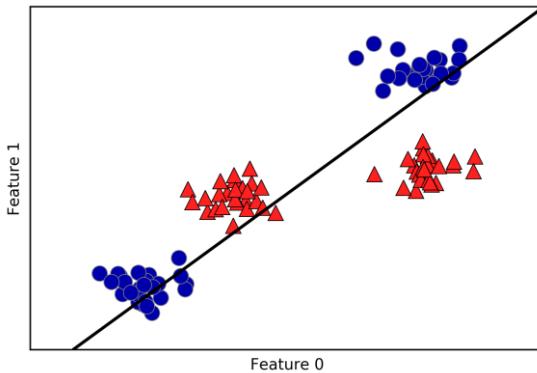


New feature space ( $x, x^2$ )

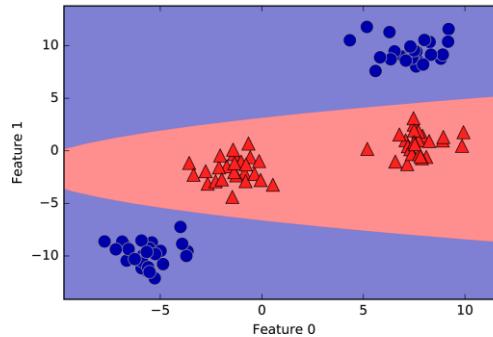


# Accommodating Nonlinearity

- Inseparable instances in 2D space becomes linearly separable if instances are mapped in a feature space of higher dimension (3D).



- What does the linear hyperplane in 3D look like in 2D? It looks non-linear.



# Linear Classifiers: Multi-Class Classification

- ▶ **One-vs-rest** classifier for each class in multi-class task.
- ▶ Each set of parameters corresponds to a class.

```
clf = LinearSVC()
```

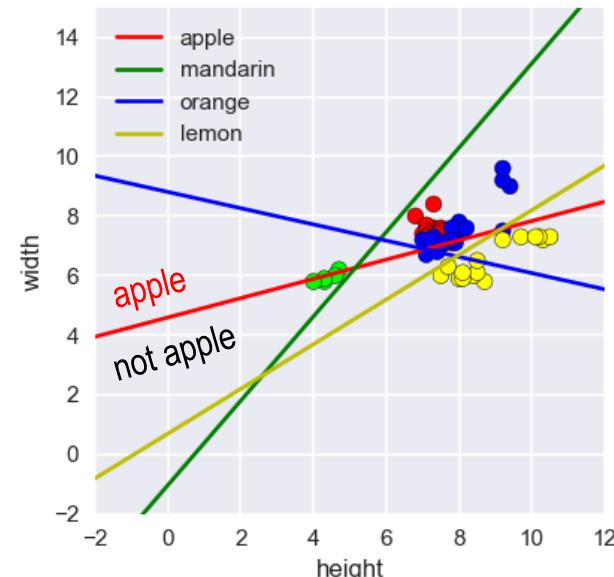
```
clf.fit(X_train, y_train)
```

```
clf.coef_
```

```
[[-0.23401135  0.72246132  
 -1.63231901  1.15222281]  
 [ 0.0849835  0.31186707]  
 [ 1.26189663 -1.68097 ]]
```

```
clf.intercept_
```

```
-3.31753728  1.19645936 -2.7468353  1.16107418]
```



$$y(\text{apple}) = -3.31753728 -0.23401135 * \text{height} + 0.72246132 * \text{width}$$

- height=2, width=6:  $\hat{y}(\text{apple}) = 0.549$  ( $\geq 0$ : predicted as apple)
- height=2, width=2:  $\hat{y}(\text{apple}) = -2.340$  ( $< 0$ : predicted as others)