


Overfitting and Its Avoidance

PF5

Learning Goals

- ▶ Generalization and overfitting
 - ▶ Overfitting and model complexity
 - ▶ Detect overfitting: fitting graphs
- ▶ Model evaluation
 - ▶ Hold-out testing
 - ▶ Cross-validation
- ▶ Overfitting avoidance
 - ▶ Nested hold-out testing
 - ▶ Grid Search with cross-validation
 - ▶ Regularization




Generalization and Overfitting

- ▶ When a model gets more **complex**, it often **fits better** to **training data**.
 - ▶ A model gets more complicated when more features are used.
 - ▶ For trees, model complexity also means the tree size (or number of nodes).
 - ▶ There is a tendency to tailor models to training data.
 - ▶ “If you torture the data long enough, it will confess”.
- ▶ **Generalization** is the goal of data mining: can this model perform well on the future **unseen data**?
 - ▶ If not, you are facing an **overfitting** problem: i.e., finding patterns that well captures training but not the unseen data.
 - ▶ Historical performance is no guarantee of future success.
- ▶ **Tradeoff** between model fit and simplicity (generalization).
 - ▶ Increasing model complexity often means a better fit to training data, but its generalization performance not always increase.
 - ▶ How to find the balance?

More Complexity, Better Generalization?


- ▶ **Holdout testing:** hold out some data, for which the target values are known, for **model evaluation** only.
- ▶ **Fitting graph:** plot **model performance**, on training and test data, against the **model complexity**.
 - ▶ Classification model performance is usually measured by prediction accuracy.




- ▶ When **model complexity increases**:
 - ▶ Training performance usually increases.
 - ▶ Generalization performance increases first, then drop(**overfitting starts**) at certain point.
 - ▶ **Sweet spot** is the highest point on the generalization performance.

Why Overfitting Occurs


- ▶ As a model gets more complex, it may pick up harmful **spurious correlations** between features and target in training data.
 - ▶ Each training set is a **finite sample** of the population, it has **variations** even without sampling bias.
 - ▶ A training set is different from the population (or the other training set).
 - ▶ The correlation between features and target in the training set may not exist in the population or future data.



Under-fitting






Good



Over-fitting

Example: Overfitting in Linear Classifiers


- ▶ Some models are by nature more likely to overfit than others.
 - ▶ **Logistic regression** is very sensitive to new instances, which can be **an outlier** that doesn't exist in future data.
 - ▶ Therefore, it tends to overfit.
- ▶ **Support vector machine** opts for **larger margin** around the decision boundary, rather than perfect separation of training instances.
 - ▶ Therefore, it is more robust and can avoid overfitting.



Overfitting Avoidance & Complexity Control

▶ Nested Holdout Testing

- ▶ Split the data into **train** & **test** fold (usually shuffle data first).
- ▶ Further split the **train** fold into two parts: **sub-train** & **validation**.
- ▶ Train models (with different complexity levels) on **sub-train**, compare their generalization performance on **validation**.
 - ▶ The best complexity level (model) yields highest generalization performance.
- ▶ With the best complexity chosen, build a model on the entire **train** fold, then evaluate its generalization performance on the **test** fold.




From Holdout Evaluation to Cross Validation

- ▶ **Holdout testing** is evaluation in laboratory setting.
 - ▶ Holdout data gives an estimate of the generalization performance; however, it is just **one single estimate**.
 - ▶ Can it be just a single particularly **lucky** (or unlucky) choice of training and test data?
- ▶ **Cross-validation** is a procedure that systematically swap out data samples for **evaluation**.
 - ▶ Use limited data to assess **confidence** in model performance estimate (e.g., mean, standard deviation).
 - ▶ Then we have a **more reliable estimate** of model performance.



(k -fold) Cross-Validation

- ▶ Split the data into k folds (usually shuffle data first).
- ▶ Iterate over data k times, in each iteration (split):
 - ▶ A different fold is held out as the **test** fold;
 - ▶ The remaining $k-1$ folds are combined as the **training** folds.
- ▶ Compute **mean**, **standard deviation** of test performances for all k folds.
 - ▶ The **mean cv score** tells the average generalization performance of a model.
 - ▶ The **standard deviation** tells how stable the generalization performance are.



Example: Churn Data

- ▶ 10-fold **Cross-validation**:
 - ▶ **Decision tree**: mean generalization accuracy is 68.6% ($sd = 1.1$)
 - ▶ **Logistic regression**: mean generalization accuracy is 64.1% ($sd = 1.3$)




- ▶ **Conclusion**: decision tree is better than logistic regression due to better average performance and stability.
 - ▶ Performance = generalization performance.

Overfitting Avoidance & Complexity Control

► GridSearch with Cross-validation

- ▶ Split the data into **train** & **test** fold (usually shuffle data first).
- ▶ Conduct ***k*-fold cross validation** on **train** fold (further split into **sub-train** and **validation**) to find the best complexity level.
 - ▶ For each complexity level, train model on **sub-train** and test it on **validation** for ***k*** times, calculate its **mean generalization performance** (cv score).
 - ▶ Compare **mean generalization performances** to find the best complexity.
- ▶ With the best complexity chosen, build a model on the entire **train** fold, then evaluate its generalization performance on the **test** fold.




If we compare 2 complexity levels with GridSearch with 5-CV, how many models have been built?

Optional: Nested CV for GridSearch

▶ Nested Cross-validation for GridSearch

- ▶ **Outer CV**: shuffles and splits the data into k -folds, and iterate over the data k times.
- ▶ **Inner CV**: in each iteration, (1) conduct k -fold cross-validation on **train** fold to find the best complexity; and (2) train a model (with best complexity) on the entire **train** fold, evaluate it on the **test** fold.
 - ▶ Each **train** fold was further split into **sub-train** vs. **validation**.
 - ▶ Each iteration may return different best complexity, as different **train** data were used in inner CV. If so, choose the one rated as the best **most of the time**.



If we compare 2 complexity levels with Nested GridSearch with 5-CV, how many models have been built?

Recap: Objective Functions

- ▶ Linear Regression:
 - ▶ minimize sum of **squared error**

$$\min_w \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

- ▶ Logistic Regression:
 - ▶ minimize sum of **log loss**

$$\min_w \sum_{i=1}^n -y_i \log(\hat{y}_i) - (1 - y_i) \log(1 - \hat{y}_i) \quad \text{where } \hat{y}_i = P(+)$$

- ▶ SVM:
 - ▶ minimize sum of **hinge loss**
 - ▶ 0 if on the correct side of its margin (and hyperplane).
 - ▶ Positive if on the wrong side of margin (or hyperplane).

$$\min_w \sum_{i=1}^n \xi_i \quad \text{where } \xi_i = \max_i [0, 1 - y_i \hat{y}_i]$$

- ▶ All **objective functions** aim to **maximize the model fit**, i.e., minimize the error (or loss) in different formats.

$$\min_w \text{error}(x, w)$$

Avoid Overfitting for Mathematical Functions

- ▶ How can we find the balance between model **fit** and **simplicity** through **objective functions**?
 - ▶ Keeping a model simple helps with generalization performance.

$$\min_w [\text{error}(x, w) + \text{penalty} (w)]$$

- ▶ **Regularization**: instead of just maximizing model fit, we optimize a combination of **fit** and **simplicity** in the objective function.
 - ▶ By minimizing the error/loss on predictions, it maximizes model fit.
 - ▶ By minimizing the penalty on coefficients, it controls the magnitude of **coefficients** to keep a model simple.
 - ▶ Absolute coefficient = the **importance of features** in predicting the target.

Regularization with (Hyper)parameters

- ▶ L2 penalty (**Ridge**): the sum of squared coefficients
 - ▶ A model has larger (absolute) coefficients brings larger penalty in the objective function.

$$\sum_{j=1}^p w_j^2$$

- ▶ L1- penalty (**Lasso**): sum of the absolute coefficients
 - ▶ Zero out many coefficients, therefore, automatically perform **feature selection** and **complexity control**.

$$\sum_{j=1}^p |w_j|$$

- ▶ **Ridge** or **Lasso regression** are regularized regression models: they minimize both residual (error) and the penalty on coefficients.
 - ▶ Linear regression is an unregularized model, which only minimize RSS.
- ▶ **GridSearchCV** is used to select best regularization (hyper)parameter α .
 - ▶ Larger α = stronger regularization: i.e., a simpler model with lots of coefficients either very small or 0.

$$\min_w [\text{error}(x, w) + \alpha * \text{penalty}(w)]$$

Regularization: Logistic Regression & SVM

► SVM

$$\min_w \left\{ \sum_{j=1}^p w_j^2 + C \sum_{i=1}^n \xi_i \right\} \text{ where } \xi_i = \max_i [0, 1 - y_i \hat{y}_i]$$

► Logistic Regression

$$\min_w \left\{ \sum_{j=1}^p w_j^2 + C \sum_{i=1}^n -y_i \log(\hat{y}_i) - (1 - y_i) \log(1 - \hat{y}_i) \right\}$$

- C is also a regularization (hyper)parameter:
 - C (always positive) is inversely proportional to regularization strength: a bigger C means weaker control of coefficient size.
- Bigger C: less tolerant of the loss --> more complicated model (better fit).
 - Larger coefficients: in SVM it also means a narrower margin with less SVs.
- Smaller C: more tolerant of the loss --> simpler model (worse fit).
 - Smaller coefficients: in SVM it also means a wider margin with more SVs.